

COM S 352 Final Exam Study Guide Fall 2022

Note: Final Exam is cumulative exam. This study guide is only for the second half semester. Please refer to the midterm exam study guide for the materials of the first half semester.

List of terms from second half reading/lecture

- **Deadlock**
 - Dining Philosophers problem
 - Conditions for deadlock
 - Mutual exclusion
 - Hold-and-wait
 - No preemption
 - Circular wait
 - Total ordering
 - Deadlock avoidance
 - Detect and recover
- **File Systems**
 - I/O Computer Architecture
 - CPU
 - Memory
 - Bus(es)
 - Peripheral/Device
 - Device Controller
 - Device Driver
 - Interrupt
 - DMA (Direct Memory Access)
 - Hard Disk Drive (HDD)
 - Sectors
 - Track
 - Rotational delay
 - Seek time
 - Random access
 - Sequential access
 - Shortest seek time first
 - File
 - Directory
 - Root directory
 - Absolute path
 - File descriptor

- System calls
 - open()
 - read()
 - write()
 - close()
 - lseek()
- hard link
- soft link
- metadata region (simple file system)
- data region (simple file system)
- contiguous allocation
- block allocation
- inode
- data bitmap
- inode bitmap
- direct indexing
- indirect indexing (single, double and triple)
- multiple-level indexing
- caching (file system blocks)
- write buffering
- Berkeley Fast File System (FFS)
- Block groups
- Path locality
- Fsync
- Journaling (write-ahead logging)
 - Metadata
 - Transaction
 - Commit
 - Checkpoint
- Flash-based SSDs
 - Cell, page, and block
 - Read page
 - Erase block
 - Program page
 - Flash translation layer (FTL)
 - Log-structured FTL
 - Garbage collection
 - Wear leveling
- Virtual Machines
 - Virtual Machine Monitor (VMM) or hypervisor
 - Supervisor mode
 - VMM page table - “physical memory” vs machine memory
 - Information gap between host OS and VMM
 - Para-virtualization

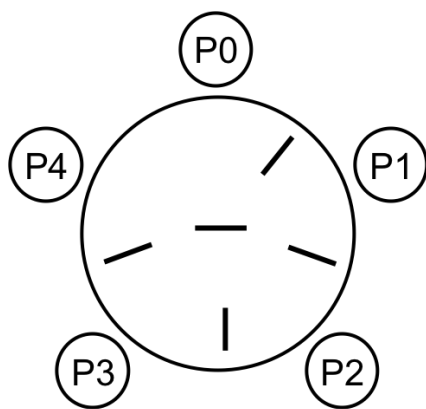
- Security
 - Confidentiality
 - Integrity
 - Availability
 - Authenticity
 - Intruder
 - Vulnerability
 - Threat
 - Attack
 - Malware
 - Trojan horse
 - Spyware
 - Ransomware
 - Sniffing
 - Masquerading
 - Man-in-the-middle
 - Denial of service (DoS)
 - Access Control
 - Object (access control)
 - Subject (access control)
 - Access mode
 - Access Control Matrix
 - Access Control List
 - Capability
 - Encryption
 - Decryption
 - Plaintext
 - Ciphertext
 - Symmetric (key system)
 - Asymmetric (key system)
 - Public key
 - Private key
 - Authentication
 - Nonce
 - Defense in Depth (design principle)
 - Layered Architecture
 - Trust
 - Trusted Computing Base (TCB)
 - Trusted Platform Module (TPM)
 - Virtual Machine
 - Sandbox
 - Honeypot
- Networking
 - Protocol

- Link
- Packet
- Switch
- Host
- Router
- Client/server
- OSI Architecture and the layers
- Internet Architecture and the layers
- Socket
- IP Address
- Port
- TCP
- UDP
- Bind, listen and accept (what a server does)
- connect (what a client does)
- bandwidth
- latency
- propagation
- Distributed Systems
 - Transparency (design principle of distributed systems)
 - Stateful (protocol)
 - Stateless (protocol)
 - NFS
 - File handle (NFS)
 - NFS commands
 - LOOKUP
 - READ
 - WRITE
 - GETATTR
 - Generation Number
 - Retry on failure
 - Idempotency
 - Cache consistency problem
 - Update visibility
 - Stale cache
 - Flush-on-close (cache)
 - Attribute cache
 - AFS
 - Scalability
 - AFS commands (only need to know these three)
 - TestAuth
 - Fetch
 - Store
 - Whole-file caching

- Volumes
- Callback
- Heartbeat protocol

Module 8

1. Consider a modified version of the dining philosophers problem, where one of the chopsticks is placed at the center of the table, available for any of the philosophers to pick up with either hand.

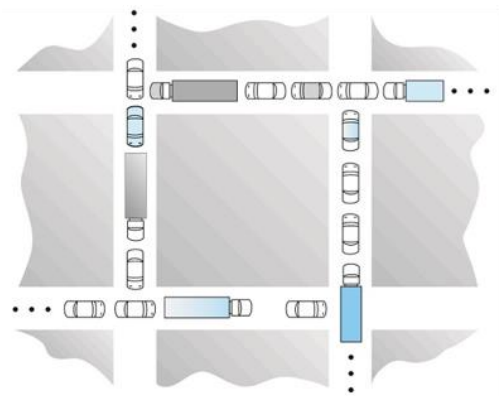


Can deadlock be prevented if no philosopher is allowed to pick up the center chopstick as their first chopstick? Explain your answer.

Yes, deadlock can be prevented.

Consider the four outer chopsticks, which are the only ones that are allowed to be picked up first. Suppose that four of the philosophers each pick up one of them and a fifth philosopher has none. Is it possible for there to be a circular wait on the chopsticks? The fifth philosopher that has none is not allowed to pick up the center chopstick first. Therefore, one of the four philosophers will pick up the center chopstick and be able to eat. In other words, because there must always be one philosopher that has no chopsticks, it is not possible for there to be a circular wait, and likewise not possible for there to be a deadlock.

2. Consider the traffic deadlock depicted in the figure below.



- a. Show that the four necessary conditions for deadlock hold in this example.
- b. State a simple rule for avoiding deadlocks in this system.

a.

Mutual exclusion – two vehicles cannot physically be in the same place at the same time

Hold and wait – There are vehicles that have blocked an intersection, thus preventing other traffic from entering the intersection (e.g., holding a resource). At the same time these vehicles are waiting for the intersection ahead of them to open up before they can clear their current intersection (e.g., waiting for a resource).

No preemption – It appears that none of the vehicles are willing or able to divert from their current path in order to open up one of the intersections.

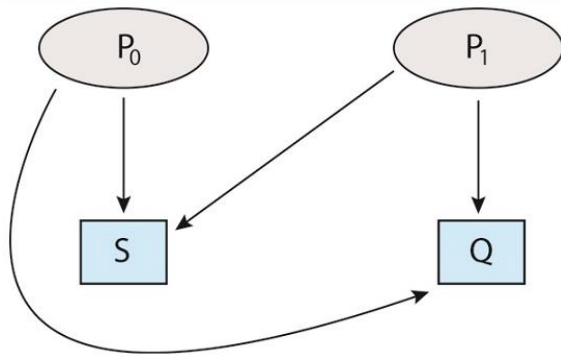
Circular wait – there are four flows of traffic and four intersections, each flow of traffic is waiting on an intersection to clear that another flow of traffic is currently occupying forming a cycle of dependencies.

b.

A vehicle may enter an intersection only if both that intersection and the intersection ahead are clear. This rule means that if three intersections are occupied by three flows of traffic, a fourth flow of traffic cannot enter the fourth intersection which prevents circular wait.

3. The resource-allocation graph below shows that both processes P_0 and P_1 need both resources S and Q to complete.

- a. Explain how a deadlock is possible in this system.
- b. Although a deadlock is possible, the system will not always enter a deadlock, the scheduler plays a role. Assume P_0 and P_1 have CPU bursts of 20 ms each, draw example Gantt charts that demonstrate FCFS and RR (with quanta=5). Use the charts to explain which algorithm is more likely to result in a deadlock?



FCFS:

Because FCFS doesn't use process preemption (different than resource preemption), it will not cause a hold and wait, therefore it will not cause a deadlock. For example, both processes in the chart below are to complete.

```
| P0 lock(s); lock(q); unlock(s); unlock(q); | P1 lock(q); lock(s); unlock(s); unlock(q); |
```

RR:

An RR scheduler uses process preemption, which can cause a deadlock, assuming there is no resource preemption. In the example below both P0 and P1 become deadlocked (enter a permanent wait state) at the statements highlighted.

```
| P0 lock(s); | P1 lock(q); | P0 lock(q); | P1 lock(s); |
```

Module 9

1. Suppose that a computer can read or write a memory word in 10 nsec. Also suppose that when an interrupt occurs, all 32 CPU registers, plus the program counter are pushed onto the stack. What is the maximum number of interrupts per second this machine can process.

To service an interrupt the CPU registers need to be written to memory and before the interrupt returns the registers need to be read back from memory. Assuming there is no caching 10 nsec is taken for each read or write of each word.

$$33 \text{ words} * 10 \text{ nsec} * 2 = 660 \text{ nsec}$$

The maximum number of interrupts per second is: $1 / 660 \text{ nsec} = 1.5 \text{ million}$

2. Consider a 320 GB SATA drive described in the figure below. Suppose the workload is 10,000 reads to sectors randomly scattered across the disk. How long with these 10,000 requests take (total) assuming the disk services requests in FIFO order?

Size	
Form factor	2.5 inch
Capacity	320 GB
Spindel speed	5,400 RPM
Average seek time	12.0 ms
Maximum seek time	21 ms
Track-to-track seek time	2 ms
Transfer rate (surface to buffer)	850 Mbit/s (maximum)
Transfer rate (buffer to host)	3 Gbit/s
Buffer memory	8 MB

disk access time = seek time + rotation time + transfer time

Seek time:

- Each request requires a seek from a random starting track to a random ending track.
- For seek time use the average of 12ms.

Rotation time:

- Once the disk head settles on the right track, it must wait for the desired sector to rotate under it.
- The average time it will take is one half of a rotation
- $0.5 * 60s/5400RPM = 5.56ms$

Transfer time:

- The disk's surface bandwidth is at least 100 MB/s, so transferring one sector (512 bytes) takes at most 5120ns (0.00512 ms).

Total time:

- Num requests * (seek + rotation + transfer)
= $10000 * (12ms + 5.56ms + 0.00512ms)$
= $10000 * 17.56512ms$

= 175.6s

3. Consider a 320 GB SATA drive described in the previous figure. Suppose the workload is 10,000 reads to 10,000 sequential sectors on the outer-most tracks of the disk. How long will these 10,000 requests take (total) assuming the disk services requests in FIFO order?

Seek time and rotation time per request are the same as above.

Transfer time: $10,000 \text{ sectors} * 512 \text{ bytes/sector} * 1/(100 \text{ MB/s}) = 51.2\text{ms}$

So, 10,000 requests will take about $12\text{ms} + 5.56\text{ms} + 51.2\text{ms} = 68.76 \text{ ms}$ (much faster than random access)

4. Linux provides a command-line utility `rm` that has the purpose of removing a file from a directory. When executing `rm` inside of the `strace` utility we can see that `rm` makes a system call to `unlink()`. Here is an example.

```
$ strace rm test.txt
...
unlinkat(AT_FDCWD, "test.txt", 0) = 0
...
```

Explain why `rm` unlinks the file. How does unlinking result in the removal of the file? Why does Linux not have a system call such as `delete()` that `rm` could use to directly remove the file?

There may be more than one file name mapped to the inode.

If the file is deleted (inode is deleted), those other directory mappings would become invalid.

Unlinking removes the directory mapping.

Only when all files (references) to the inode are unlinked will the inode be deleted.

5. Linux provides the command `mv` that has the purpose of “moving” a file from one file path name to another file path name. When executing `mv` inside of the `strace` utility we can see that `mv` makes a system call to `rename()`.

```
$ strace mv oldname.txt newname.txt
...
renameat(AT_FDCWD, "oldname.txt", AT_FDCWD, "newname.txt") = 0
...
```

Is there any difference at all between renaming a file and just copying the file to a new file with the new name, followed by deleting the old one? Explain.

Renaming a file means modifying the directory mapping of file name to inode. The file contents do not need to be copied.

The alternative approach would not be the same, besides having an extra copy, there would be complications such as other hard links to the file becoming irrelevant.

Module 10

1. How many disk operations are needed to fetch the inode for the file `/var/log/boot.log`? Assume that the inode for the root directory is in memory, but nothing else along the path is in memory. Also assume that all directories fit in one disk block.

Assuming the inode for the root directory is in memory, the reads will be:

1. Root data block
2. Var inode block
3. Var data block
4. Log inode block
5. Log data block
6. boot.log inode block

Therefore, it will be 6 read operations.

2. In the Very Simple File System (VSFS) the inodes are kept together near the start of the disk. An alternative design is to allocate an inode when a file is created and put the inode at the start of the first data blocks of the file.

- a. Assume a hard disk drive (HDD) is being used. What is an advantage to putting inodes at the start of the disk? What is an advantage to putting inodes at the start of the first data blocks of the file?

When inodes are together at the start of the disk there may be less seek time to find multiple files in a directory. When inodes are close to the contents of the file there is less seek time to go from the inode to the data of the file.

- b. In the Fast File System (FFS) the inodes for a file are placed near the start of the same block group as the file's data blocks. Again assuming a HDD, explain an advantage to this approach rather than placing all the inodes near the start of the disk.

This has the same advantages of putting inodes at the start of the disk for files that are close together in the directory tree (thus get put in the same group). But there is an additional advantage that the inodes are closer to the file data than they would be if placed at the start of the disk.

- c. Now assume a Solid State Drive (SSD) is being used. Does the FFS approach still have any advantages or would the alternative design of putting the inode at the start of the first data blocks of the file be better?

SSD still benefits from faster sequential access, however, there is not seek latency so whether blocks are close or far in the logical address space doesn't make a difference. Therefore, there would be no difference in either of the two approaches

3. What would happen if the bitmap containing the information about the free disk blocks was completely lost due to a crash? Is there any way to recover from this disaster, or is it bye-bye disk?

It is possible to completely recover the bitmap. Starting at the root directory every directory and file can be recursively visited. When the inode of a file or directory is visited it is possible to discover what other blocks that inode is pointing to. If the bits in the bitmap are marked as used for each block encountered, at the end the bitmap will be recovered.

4. Some file systems allow disk storage to be allocated at different levels of granularity. For instance, a file system could allocate 4 KB of disk space as a single 4-KB block or as eight 512-byte blocks.

- a. How could we take advantage of this flexibility to improve performance?

The advantage is less internal fragmentation will lead to less wasted space. Suppose a file only uses half of a block, then the other half would have been wasted space. With 512-byte subblocks the file will occupy exactly two sub-blocks and waste no space.

- b. What modifications could be made to the bitmap free-space management scheme in order to support this feature?

There are many possible ideas. The most naive approach would be to make a bitmap that is 4 times larger so that every bit represents one sub-block.

A more efficient approach might be to combine the ideas of a bitmap and free list. The bitmap would indicate if the entire block is free. When no free blocks remain a linked list indicates which sub-blocks are free.

5. Consider a file system that uses Unix-style inodes (index nodes) to represent files. The system uses 16-KB blocks and 8-byte pointers. Multi-level indexing is used with an inode containing 12 direct pointers, 1 single indirect pointer, 1 double indirect pointer and 1 triple indirect pointer.

- a. What is the maximum size of a file that can be stored in the file system? Show calculations.

**A block can hold $16\text{KB} / 8 \text{ bytes} = 2,048$ pointers
 $(12 + 2048 + 2048^2 + 2048^3) * 16\text{KB}$**

$$(12 + 2^{11} + 2^{22} + 2^{33}) * 2^{14}$$

$$8,594,130,956 * 16,384$$

$$(8,594,130,956 * 16,384) / 2^{40} \text{ TB} = 128 \text{ TB}$$

- b. For each of the following addresses describe how many blocks (including the inode and data block) need to be read to perform a read of the data at the address. Explain your answers.

4,096
32,768
1,048,576
268,435,456

$$4,096 / 16,384 = \text{block \# 0}$$

The block is less than 12, therefore it is directly pointed to by the inode, so 2 reads.

$$32,768 / 16,384 = \text{block \# 2}$$

The block is less than 12, therefore it is directly pointed to by the inode, so 2 reads.

$$1,048,576 / 16,384 = \text{block \# 64}$$

The block is ≥ 12 but less than 2060, therefore it is pointed to by a single indirect pointer, so 3 reads.

$$268,435,456 / 16,384 = \text{block \# 16,384}$$

The block is $\geq 2,060$ but less than 4,196,364, therefore it is pointed to by a double indirect pointer, so 4 reads.

6. Suppose an SSD consists of raw flash with the following performance characteristics: reading one page takes $25\mu\text{s}$, programming one page takes $200\mu\text{s}$ and erasing one block takes $1500\mu\text{s}$ (note these refer to physical SSD pages and blocks). The following figure represents the current state of the SSD.

Table:	100->3	101->4	102->5					
Block:	0				1			
Page:	00	01	02	03	04	05	06	07
Content:	a1	a2	a3	b1	b2	b3		
State:	V	V	V	V	V	V	E	E

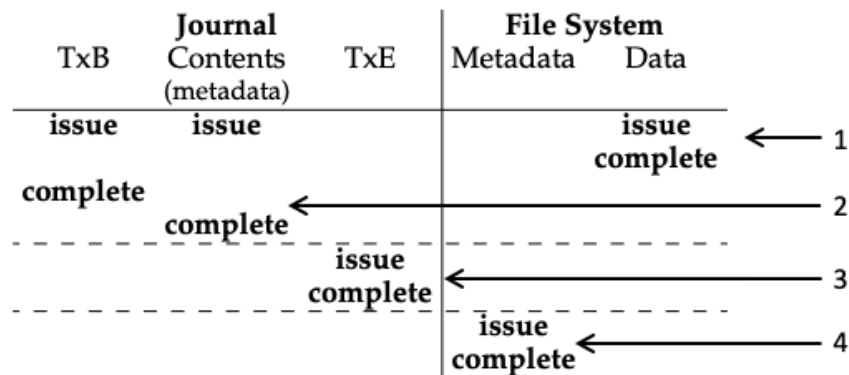
How much time will it take for garbage collection to reclaim physical block zero? Assume all operations must be performed sequentially (no parallel hardware optimization) and none of the pages are cached. Explain your result.

Read b1 from page 3
Program b1 to page 6
Erase block 0.

Therefore the time is: $25 + 200 + 1500 = 1725\mu\text{s}$

Module 11

1. The figure below shows a timeline (top to bottom) of when different actions required to write to a block to a file occur. Explain how the file system recovers from a failure at each of the 4 points indicated and say if any data is lost after the recovery.



1. No recovery as the data has not been logged to the journal or file system.
2. No recovery as the metadata has not been logged to the journal or file system.
3. No recovery as we don't know if the journal is complete (TxE has not been issued yet).
4. Yes. We can replay the journal to recover the metadata in the file system. Data is already in the file system since it is done before the issue of TxE.

2.

- a. Suppose the file system issues a write command to update the contents of virtual page 2000 with data "c1". Assume the current state of the SSD drive is shown in the figure below. Draw the most likely state of the SSD after the write has been completed.

Table:	100 → 0	101 → 1	2000 → 2	2001 → 3	Memory
Block:	0				
Page:	00 01 02 03	04 05 06 07	08 09 10 11		
Content:	a1 a2 b1 b2				Flash Chip
State:	V V V V	i i i i	i i i i		

- b. Now suppose, following the previous write, the file system deletes a file and removes virtual pages 100 and 101. Assume the FTL decides to garbage collect block 0. Draw the most likely state of the SSD after the garbage collection has been completed.

a.

table: 100->0, 101->1, 2000->4, 2001->3
| a1 | a2 | b1 | b2 | c1 |

b.

table: 2000->4, 2001->3

| e | e | e | e | c1 | b2 |

3. Suppose a guest process is running on top of a guest operating system on top of a host operating system. Consider the memory mappings below. Assume page size is 1KB (1,024 bytes).

Guest OS Page Table

VPN (Virtual Page Number)	PFN (Physical Frame Number)
0	10
1	2
2	3
3	5
4	0
5	1
6	6
7	7

Hypervisor (Shadow) Page Table

PFN (Physical Frame Number)	Machine
0	11
1	12
2	5
3	6
4	9
5	10
6	13
7	14

Where on the machine is virtual byte 2,555 stored?

$$\text{VPN} = 2,555 \% 1,024 = 2$$

$$\text{PFN} = 3$$

$$\text{Machine frame number} = 6$$

$$\text{Base address} = 6 * 1,024 = 6,144$$

$$\text{Offset} = 2,555 - (1,024 * 2) = 507$$

$$\text{Machine address} = 6,144 + 507 = 6,651$$

Module 12

1. Select the security goal from the CIA Triad (confidentiality, integrity and availability) that has been violated in each scenario.

- a. A user replaces wininit.exe on a Windows file system with a modified version that contains a backdoor to allow unauthenticated access to the machine in the future.
 - b. A botnet (a group of compromised Internet-connected devices) is being used to target a webserver with high amounts of network traffic with the objective of overwhelming the server's resources.
 - c. A user is logged into their account at Secure Bank & Trust, an attacker has launched a man-in-the-middle attack that allows the attacker to observe the communication.
-
- a. **Integrity: system resource replaced by unauthorized**
 - b. **Availability: distributed denial of service DDoS**
 - c. **Confidentiality: observing secret/private data**

2. Discuss how the asymmetric encryption algorithm can be used to achieve the following goals.

- a. Authentication: the receiver knows that only the sender could have generated the message.
- b. Secrecy: only the receiver can decrypt the message.
- c. Authentication and secrecy: only the receiver can decrypt the message, and the receiver knows that only the sender could have generated the message.

Let,

k^S_e be the public key of the sender,

k^r_e be the public key of the receiver,

k^S_d be the private key of the sender, and

k^r_d be the private key of the receiver.

- **Authentication:** Authentication is performed by having the sender send a message that is encoded using Private Key of Sender, k^S_d .
- **Secrecy:** Secrecy is ensured by having the sender encode the message using Public Key of the Receiver k^r_e .
- **Authentication and Secrecy:** Both authentication and secrecy are guaranteed by performing double encryption using both Private Key of Sender k^S_d and Public Key of Receiver k^r_e .

3. Many computer architectures offer more than just two security modes, for example, they have user, kernel and machine mode. If the operating system is trusted to execute with kernel mode privileges, why is there need for a machine mode?

Machine mode allows adding an extra level of protection such as putting the OS in a virtual machine, hypervisor, or sandbox or using a Trusted Platform Module.

4. How do operating systems that use paged virtual memory prevent processes from reading or writing the physical memory of other processes? Be specific about the mechanisms in place that prevent this.

All memory accesses by a process are treated as virtual.

The hardware MMU translates virtual addresses to physical addresses by looking in the page table.

If the page table does not have a translation, then the memory access is not allowed.

The page table only maps to addresses the process is allowed to access.

5. An operating system provides access control which prevents unauthorized reading of data, what are the advantages of encrypting data stored in the computer system?

The advantage is security in depth. Even though the OS has access control, it should not be relied on as the only security mechanism for sensitive data such as passwords. Systems are hacked, if the attacker gets access to the password file it is better if the file is encrypted to add an extra layer of defense.

6. Describe each of the following three kinds of access control mechanisms in terms of (a) ease of determining authorized access during execution, (b) ease of adding access for a new subject, (c) ease of deleting access by a subject, and (d) ease of creating a new object to which all subjects by default have access.

- a. per-subject access control list (that is, one list for each subject tells all the objects to which that subject has access)
- b. per-object access control list (that is, one list for each object tells all the subjects who have access to that object)
- c. access control matrix

a.

- i. **Ease of determining authorized access during execution – it is easy to find the object that is being accessed in the list of objects for the subject**
- ii. **Ease of adding access for a new subject – it is easy to add a new subject and add the object into that subjects list.**
- iii. **Ease of deleting access by a subject – find the object in the subjects list and remove it from the list.**

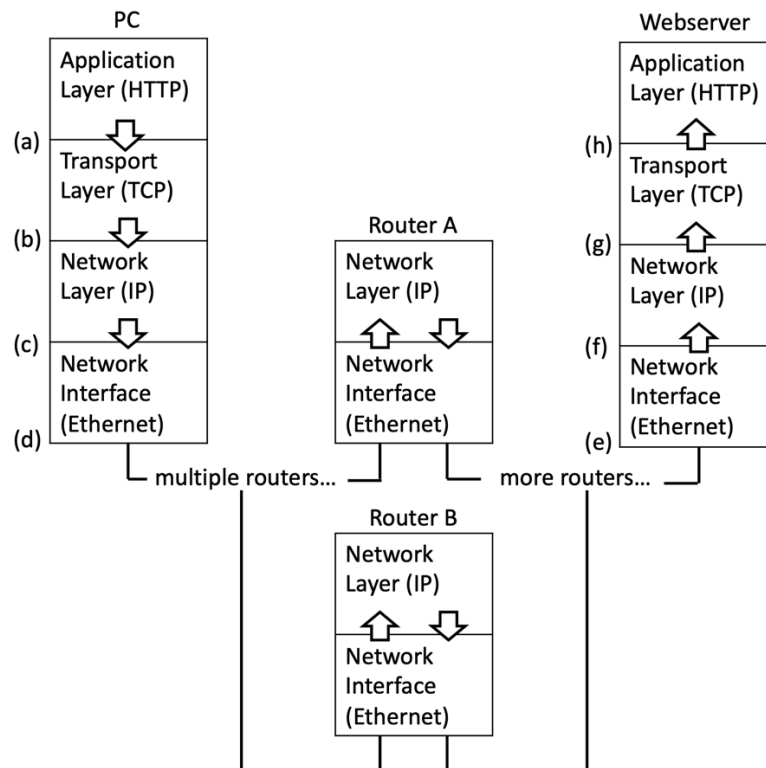
- iv. Ease of creating a new object to which all subjects by default have access – this is more time consuming, it requires visiting each subject and adding the object to that subject's list
- b.
 - i. Ease of determining authorized access during execution – it is easy to find the find the subject in the list for the object
 - ii. Ease of adding access for a new subject – it is easy to add a new subject to the list for the object
 - iii. Ease of deleting access by a subject – find the subject in the object's list and remove it from the list.
 - iv. Ease of creating a new object to which all subjects by default have access – create a new object which a list of every subject.
- c.
 - i. Ease of deterring authorized access during execution – find the row and column in the matrix
 - ii. Ease of adding access for a new subject – add a new subject row and set the value in the correct column
 - iii. Ease of deleting access by a subject – fix the row and column in the matrix and clear the access
 - iv. Ease of creating a new object to which all subjects by default have access – add an object column and the value in every row.

7. Suppose a per-subject access control list is used. Deleting an object in such a system is inconvenient because all changes must be made to the control lists of all subjects who did have access to the object. Suggest an alternative, less costly means of handling deletion.

Keep both a per-subject and per-object access control list. The per-object control list can be used to efficiently find all subjects that have access to the object. Then removing the object for all subject lists is more efficient.

Module 13

1. The diagram shows the scenario of a PC making a request of a webserver. The Internet is represented between the PC and the webserver. There are two possible paths for packets to take, one passes through router A and the other through router B. Assume that it is equally likely for packets to take either path. Use the example to answer the following questions.



a) Explain the difference between the transport layer and the network layer.

Both are part of the network stack, and the transport layer relies on the network layer to complete its tasks. The transport layer provides communication between processes (usually on different machines), the processes are identified by port numbers, and it is expected to be resilient to packet loss or out of order packets. The network layer provides communication between hosts (e.g., machines), the hosts are identified by IP addresses, and no attempt is made to resent, or reorder lost or out of order packets.

b) Why do the routers not have a transport layer?

Routers participate at the network layer and below, they look at each packet's IP address and forward the packet to the next appropriate node, they do not try to check for lost or out of order packets. Adding transport layer mechanisms at the router level would make routing slow, expensive, and inflexible.

- c) For efficient sharing of the network infrastructure, TCP breaks long application layer messages into small packets (typically less than 1500 bytes). Explain why the packets arriving at the webserver can be out of order.

The hosts (PC and webserver) have no control over the route packets will take over the network, some packets might go to router A and others to router B.

Communication over the network is asynchronous, some packets may be lost, some packets may take a longer/slower route, so packets from router A and router B could arrive at the webserver in any order.

- d) The arrows in the diagram show an example TCP communication from the PC to the Webserver, the layer boundaries have been labeled (a) to (h). List all boundaries (a to h) where packets might **arrive** out of order.

On the PC, the application must provide the correctly ordered message to the transport layer, where there would be no reason for the transport layer, network layer or network interface to change the order of packets. Once the packets reach the network there is not guarantee over which route packets will take or how long a particular route will take. The transport layer (TCP) on the webserver maintains a receive window, it collects packets in the window and uses acknowledgements to have lost packets resent. Only a complete (contiguous) window of packets is forwarded to the application layer.

2. What are the three parts of a socket? Which layer of the 7-layer OSI model does each part deal with?

Protocol – transport layer

Port – transport layer

IP Address – network layer

3. What are the advantages and disadvantages of TCP and UDP? Give examples of applications where each might be preferred.

TCP is reliable as it guarantees the delivery of data to the destination router. A disadvantage of TCP is that it is slower and more expensive because it relies on establishing and maintaining a connection which requires extra packets being sent.

UPD is a connectionless protocol, packets are sent without any guarantee of delivery. The advantage is no overhead – no handshake packets sent to establish a connection, no acknowledgement packets sent for every data packet, etc. The disadvantage is that packets can be lost or delivered out of order, it is up the application to deal with that.

TCP is common for file transfer (e.g., FTP), email (e.g., SMTP), and web (e.g., HTTP). UDP is common for multimedia applications (e.g., VoIP) and client-server communication that doesn't need to send large messages (e.g., DNS).

4. Calculate the total time required to transfer a 1000-KB file in the following cases, assuming an RTT of 50 ms, a packet size of 1 KB data:

(a) The bandwidth is 1.5 Mbps, and data packets can be sent continuously.

Bandwidth = 1.5 Mbps = 0.1875 MB/s

Packet size = 1 KB = 1,024 B

Transmit time per packet = size/bandwidth = 1,024B / (0.1875B/s * 1024*1024) = 5.200ms

Transmit time for 1,000 packets = 5,200ms

Propagation + queuing time = RTT / 2 = 25ms

Transfer time (latency) = 5,200ms + 25ms = 5,225ms

(b) The bandwidth is 1.5 Mbps, but after we finish sending each data packet we must wait one RTT (for data receive and acknowledgement) before sending the next.

(5.200ms + 50ms) * 1,000 = 55,500ms

Module 14

1. Suppose a distributed system is using acknowledgement-based communication:
 - a. How does the sender distinguish between the message being lost, the receiver being not available, or the acknowledgement being lost?

The sender has no way to know.

- b. Explain, how the sender deals with a lost acknowledgement.

After a timeout, it resends the original message. Messages have a sequence number attached so the receiver will know if it receives the same message twice.

- c. Explain, how the sender deals with an acknowledgement that is significantly delayed.

From the perspective of the sender this is the same as a lost acknowledgement. If the acknowledgement does not arrive before the timeout, it resends the original message.

2. In most operating systems file access permissions are checked only when the file is opened and not with each read or write. Thus if you open a file successfully for both reading and writing, obtaining file descriptor f , and if you subsequently change the file's access permissions to read-only, you can still write to the file using f .

- a. Explain why supporting this feature is difficult in distributed file systems whose servers do not hold open-file state, such as NFS.

A stateless server such as NFS does not remember what it told a client in a previous response. Because of this, for security, the NFS server must check file permission on every request.

- b. How might you "approximate" this feature so that it's supported on NFS for what's probably the most common case: the file's owner is accessing the file?

If the file's owner is always given both read and write permission at the server (which seems reasonable given that it is the owner), then additional permission restrictions can be supported at the client.

3. Suppose we wish to add to basic NFS a remote procedure for appending data to the end of a file.

- a. Why does NFS not already have an *APPEND* procedure?

Append is not an idempotent operation, so it was omitted from the client-server protocol.

- b. Explain why such a remote procedure is needed: why can't its effect be obtained through the use of the *GETATTR* procedure (which, among other things, returns the size of a file) and the *WRITE* procedure?

The approach of using GETATTR followed by WRITE is not an atomic operation. This could result in a race condition between client where they overwrite each other appends.

- c. Assuming a well behaved network and a server that never crashes, will there be any problems if an append request or response is lost? Explain.

Yes, the client will assume the request failed and resend to the append. This can potentially result in multiple appends to the file.