

# Recap

Process data structures:

- User space
- PCB (process control block)

Multiprogramming:

- Scheduler, Process states/queues
- Context switch

Process API: fork, wait, exec, pipe, dup

xv6-riscv

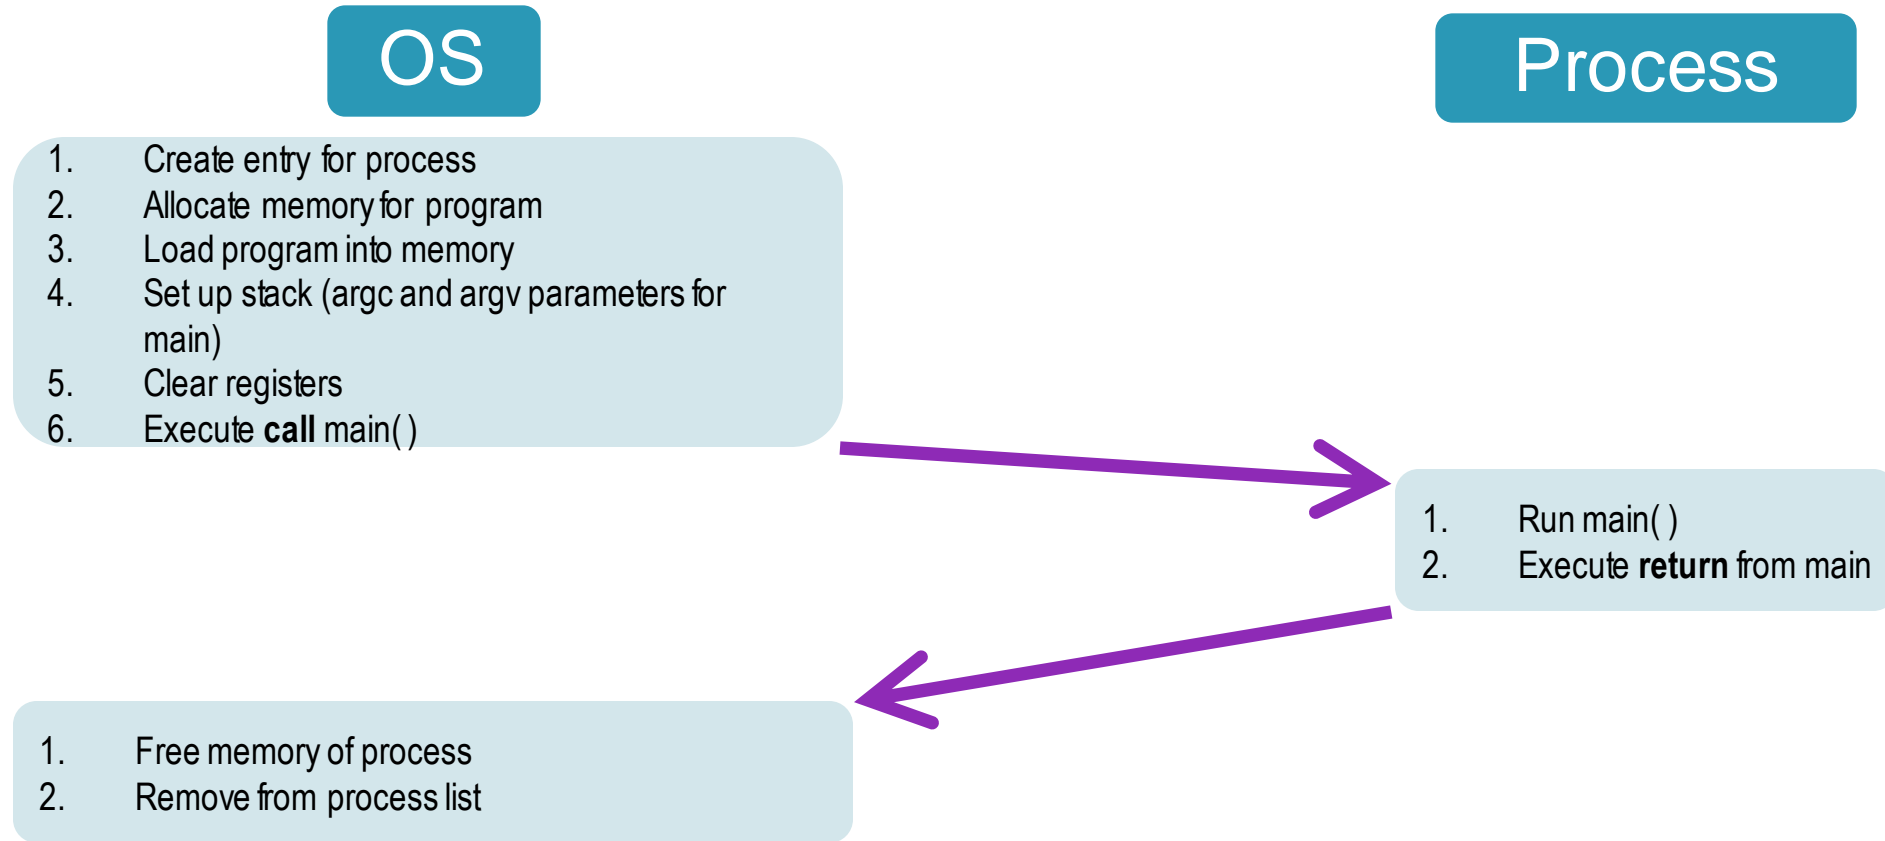
# L3: Limited Direct Execution

(Based on Chapter 6)

Spring 2023

How to run processes?

# Direct Execution (without limits)



## What is wrong?

- Process has unrestricted access to memory and other resources.
- OS has no way to switch another process, must wait for program to finish.

How to restrict process' operations?

# Restricting Access with Processor Mode

CPU has bit that indicates if in *user mode* or *kernel mode*

When in user mode, *“privileged” instructions* not allowed and memory boundaries are enforced

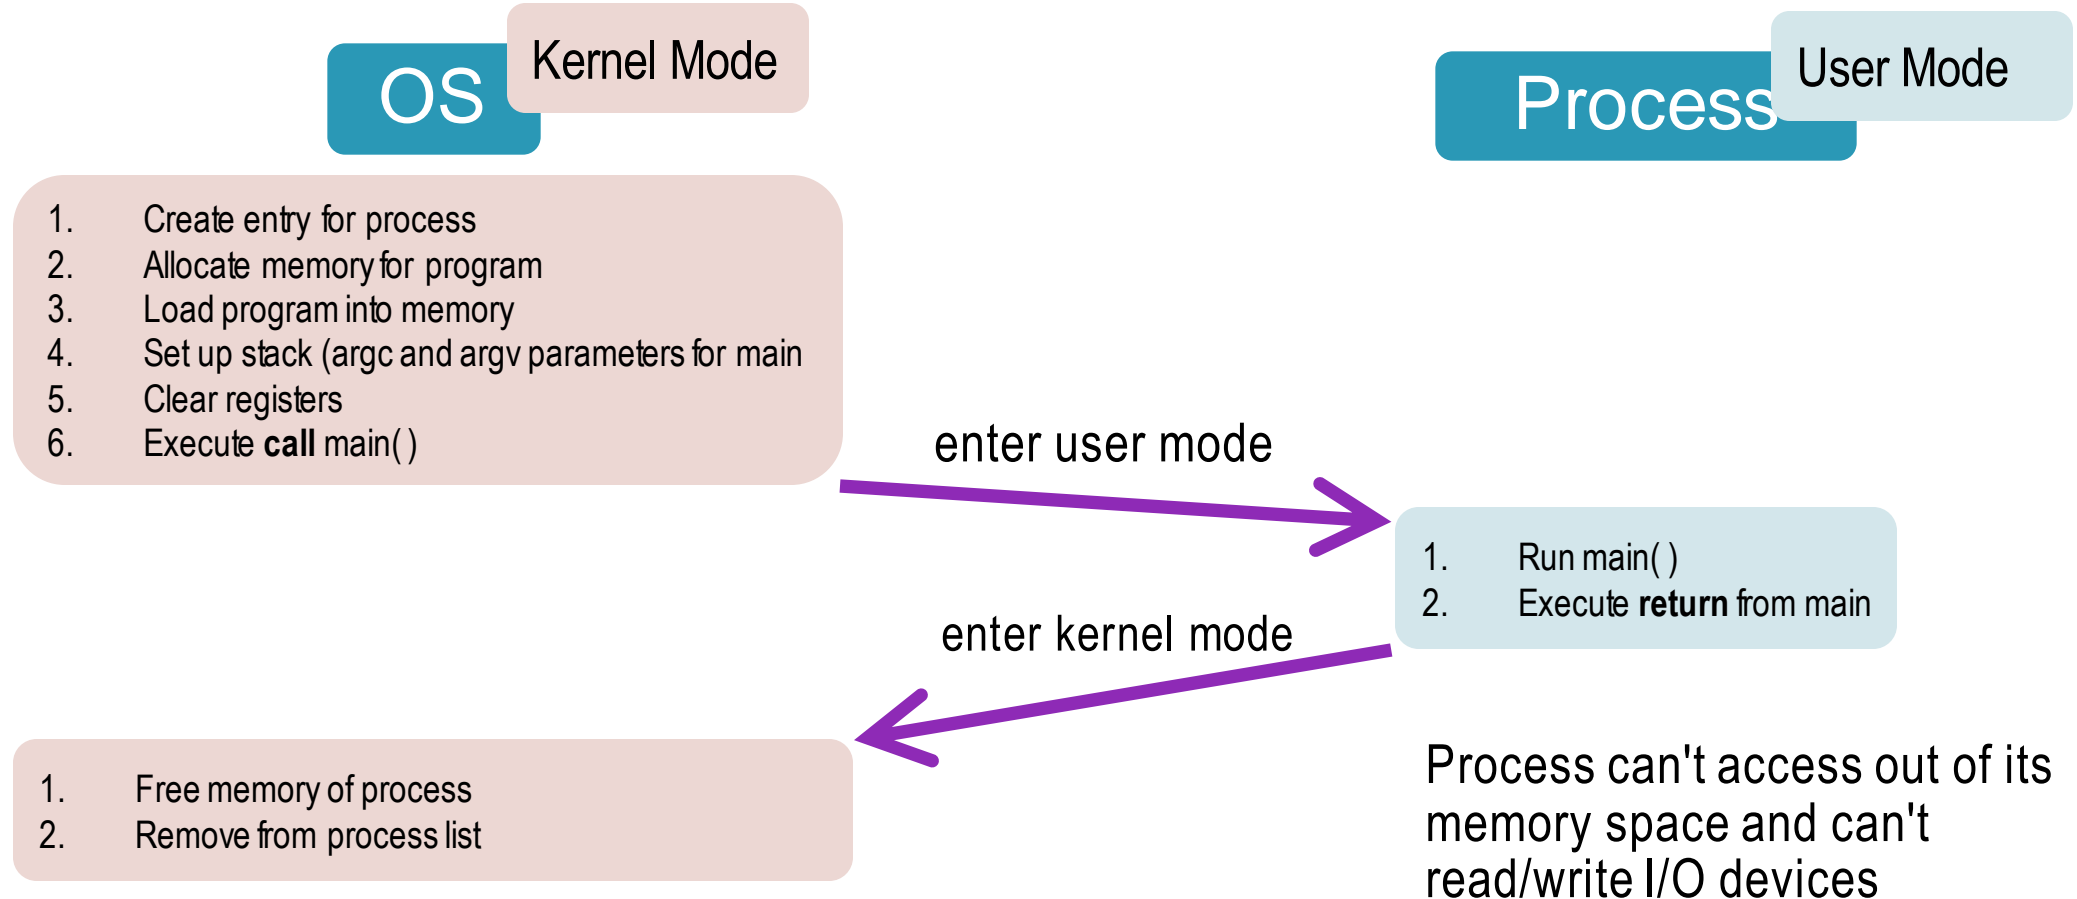
- Cannot read/write outside of address space bounds

- Cannot read/write I/O devices

When in kernel mode, all instructions are allowed

User processes only execute in user mode, the OS executes in kernel mode

# Processor Mode



# System Call

Problem: How can a process perform privileged operations, for example read from a file? *System call!*

If no different modes, user could just make a functional call to "kernel" libraries.

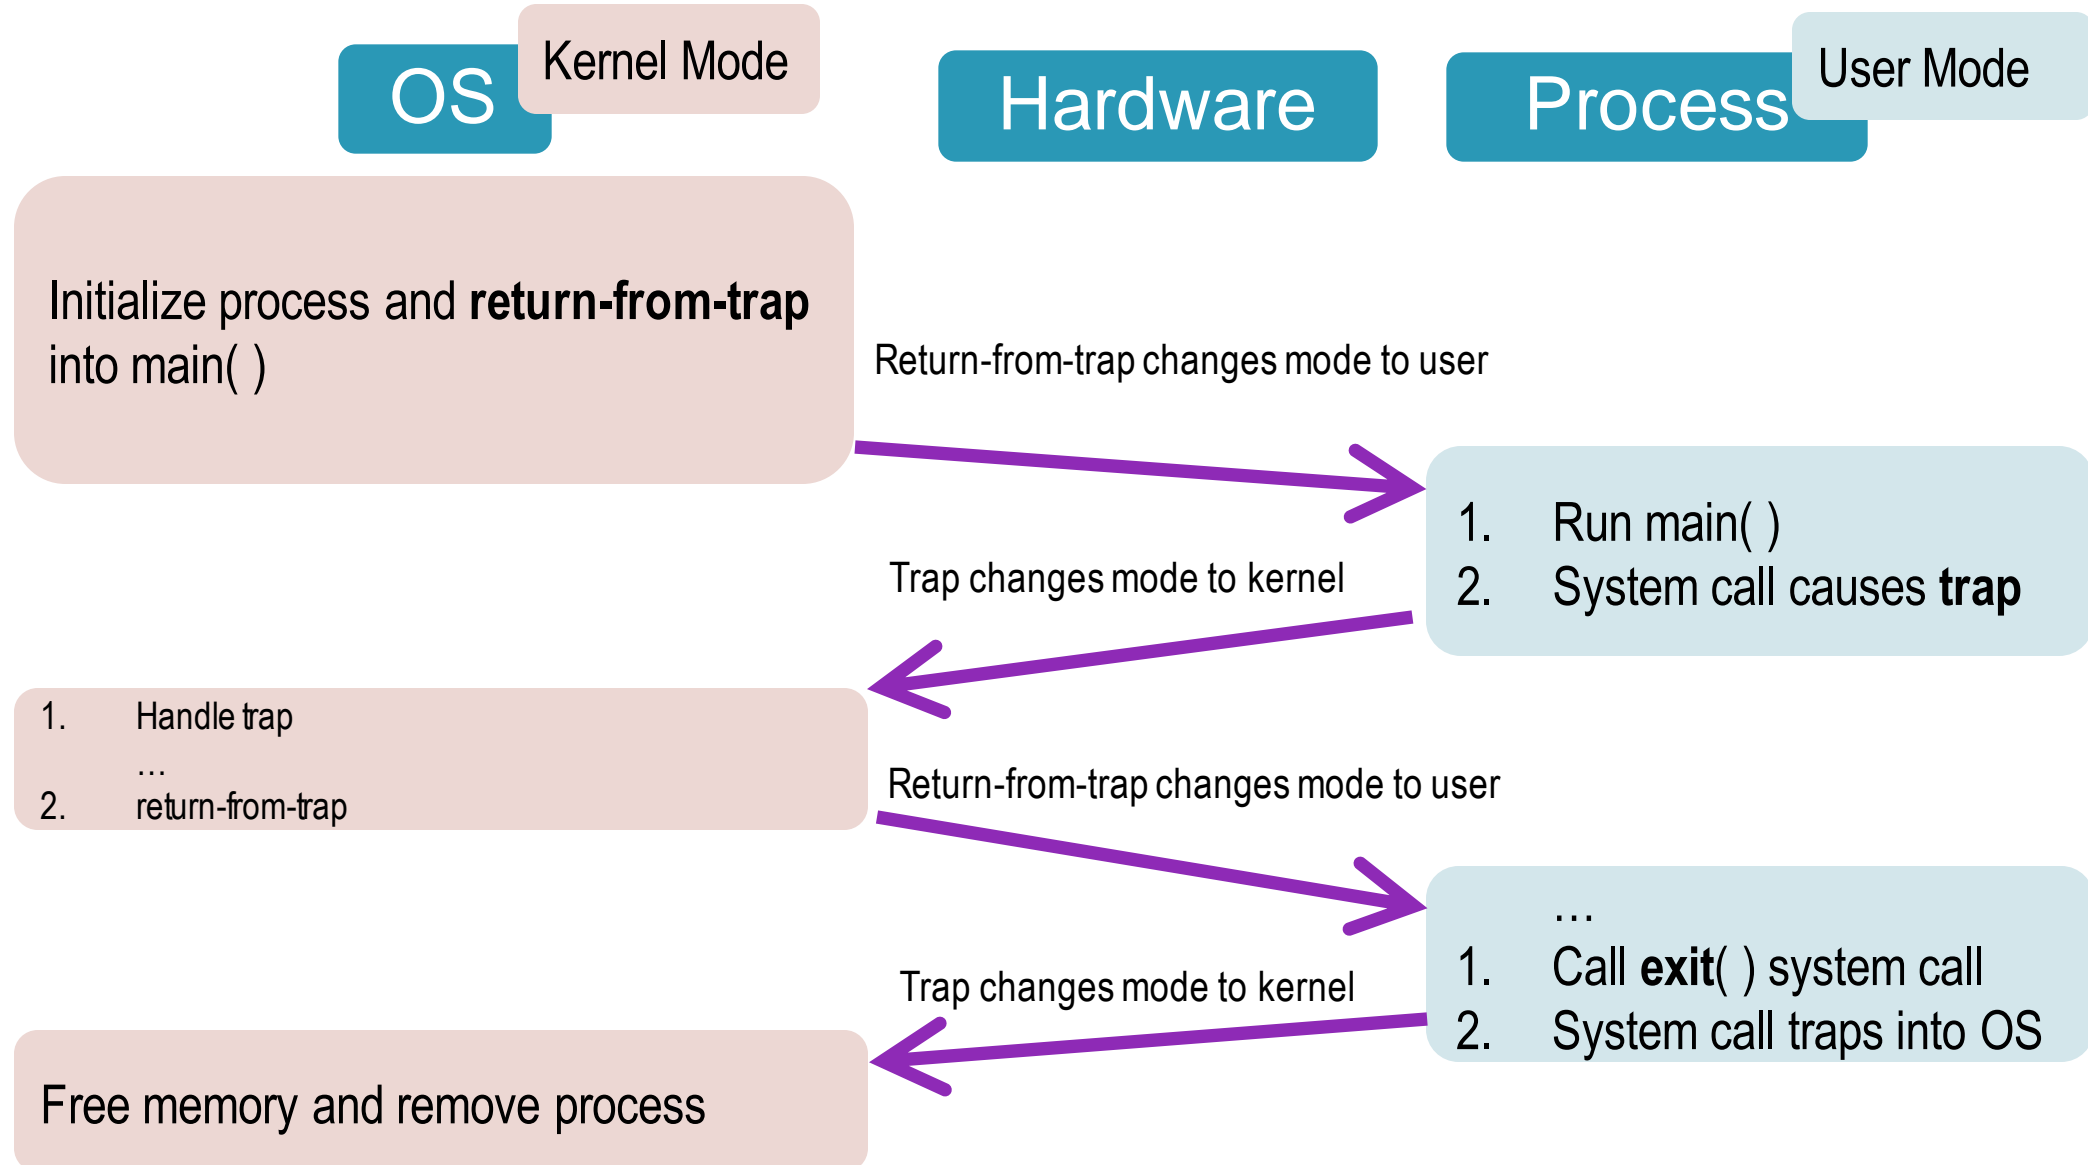
A *system call* differs from a normal function call: it transfers control to the OS.

A user process initiates a system call by causing a *trap* (a software generated interrupt)

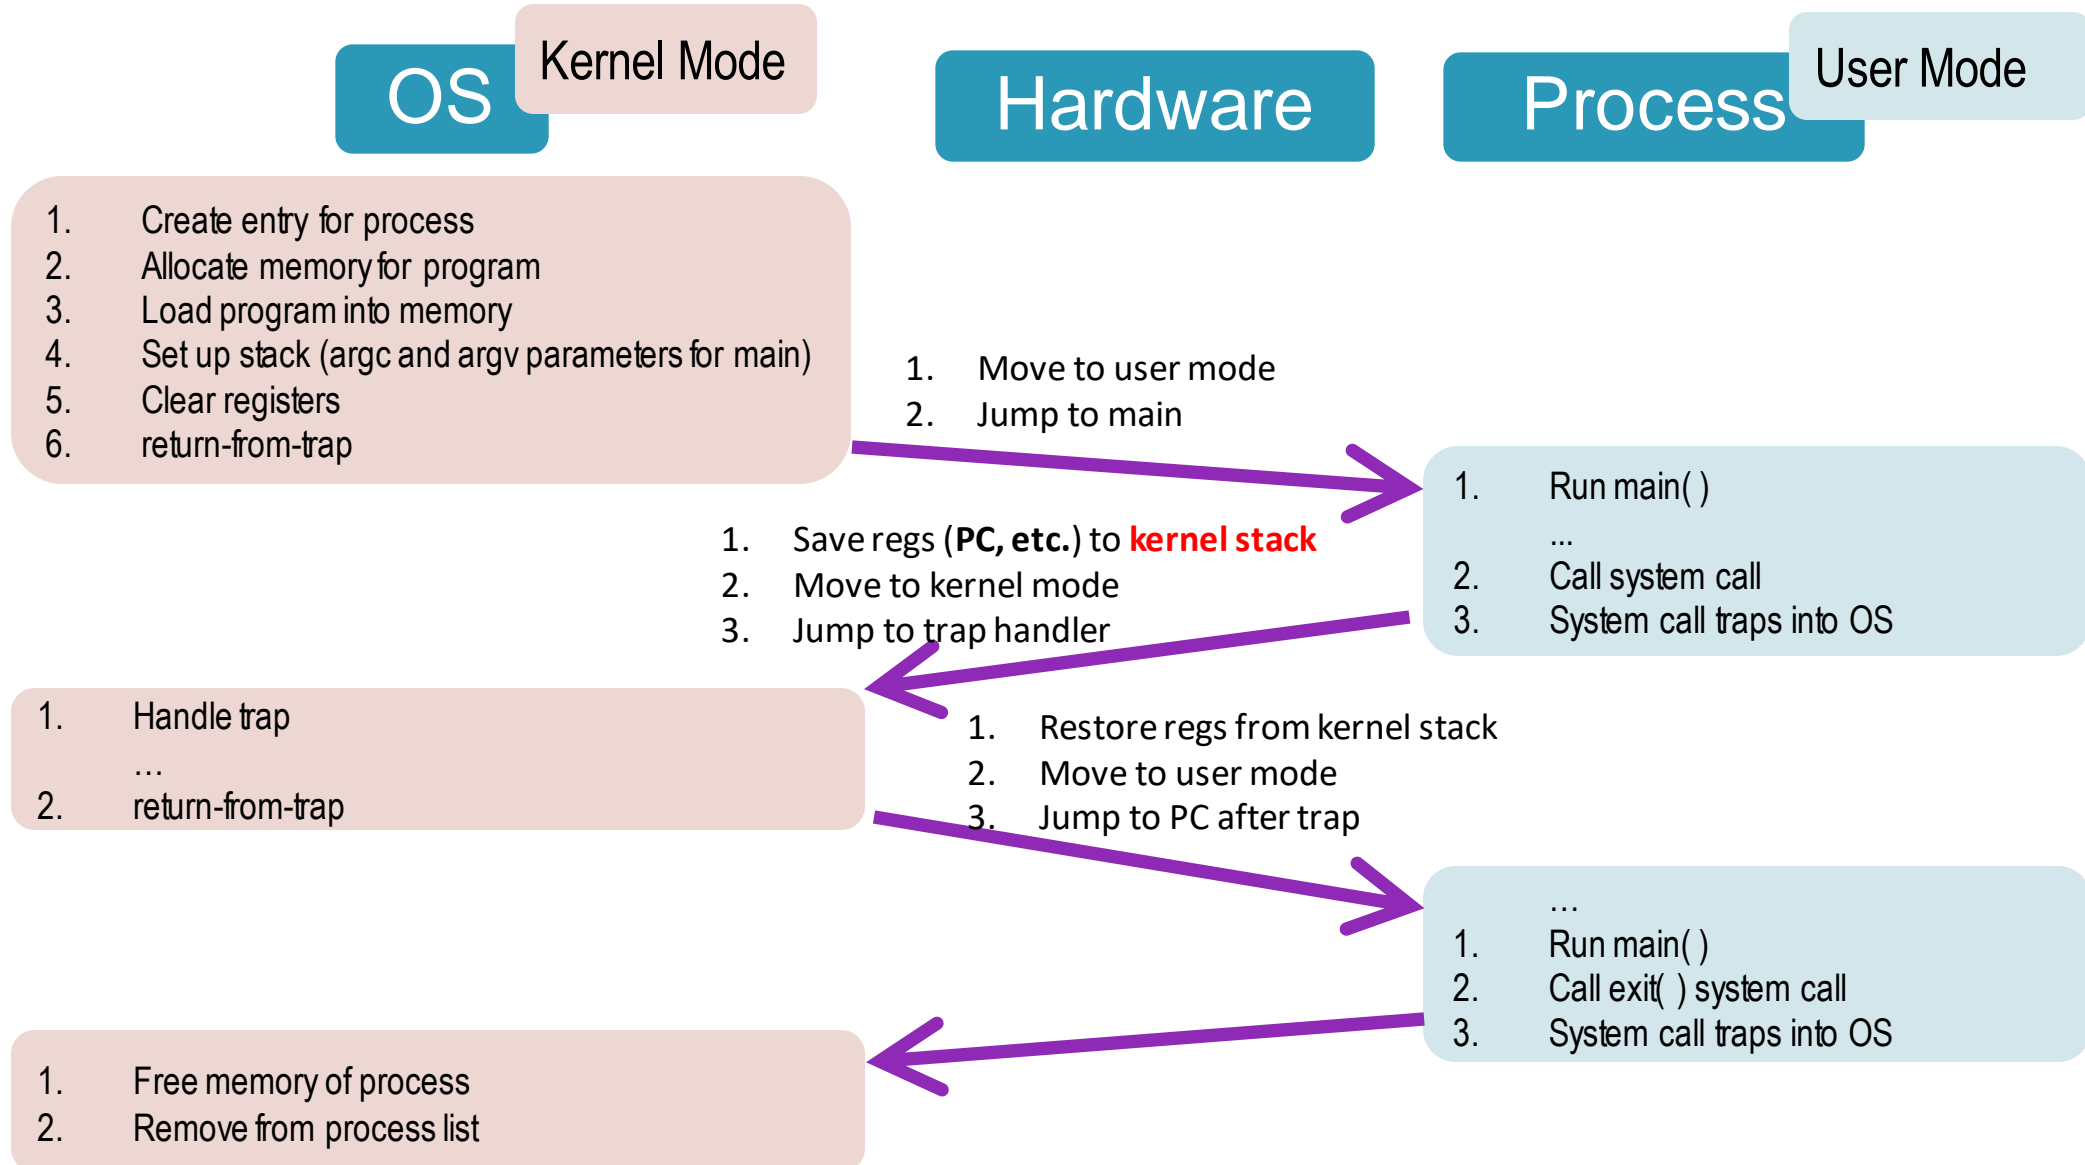
- For example, the RISC V instruction that causes a trap is *ecall*



# System Call



# System Call Details



# Process Control Block (struct proc) in xv6

```
proc.h

// Saved registers for kernel context switches.
struct context {
    uint64 ra;
    uint64 sp;

    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;
};

enum procstate { UNUSED, USED, SLEEPING,
RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum procstate state;           // Process state
    void *chan;                     // If non-zero, sleeping on channel
    int killed;                     // If non-zero, have been killed
    int xstate;                     // Exit status to be returned to parent's wait
    int pid;                        // Process ID

    // proc_tree_lock must be held when using this:
    struct proc *parent;            // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;                  // Virtual address of kernel stack
    uint64 sz;                      // Size of process memory (bytes)
    pagetable_t pagetable;          // User page table
    struct trapframe *trapframe;    // data page for trampoline.S
    struct context context;         // swch() here to run process
    struct file *ofile[NOFILE];    // Open files
    struct inode *cwd;              // Current directory
    char name[16];                  // Process name (debugging)
};
```

# System Call Details: trap table

OS

Kernel Mode


Hardware

Initialize trap table (map:  
interrupt number -> address of  
interrupt handler)

Remember address of syscall handler

# System Call Example

Library functions like `printf` typically have some code that executes in user mode, it then makes one or more system calls.



How to switch between process?

More fundamentally, how the OS can **regain the control** after it starts a process?

# A Cooperative Approach: Wait for System Calls

A well-behaving process makes a *yield* system call after it has executed for a while.

A process initializes an I/O operation by making a system call (e.g., *read*, *write*).

Once trap to OS: scheduler can switch process.

# A Non-cooperative Approach: Timer Interrupt

OS

Kernel Mode

Hardware

Initialize trap table (map:  
interrupt number ==> address of  
interrupt handler)

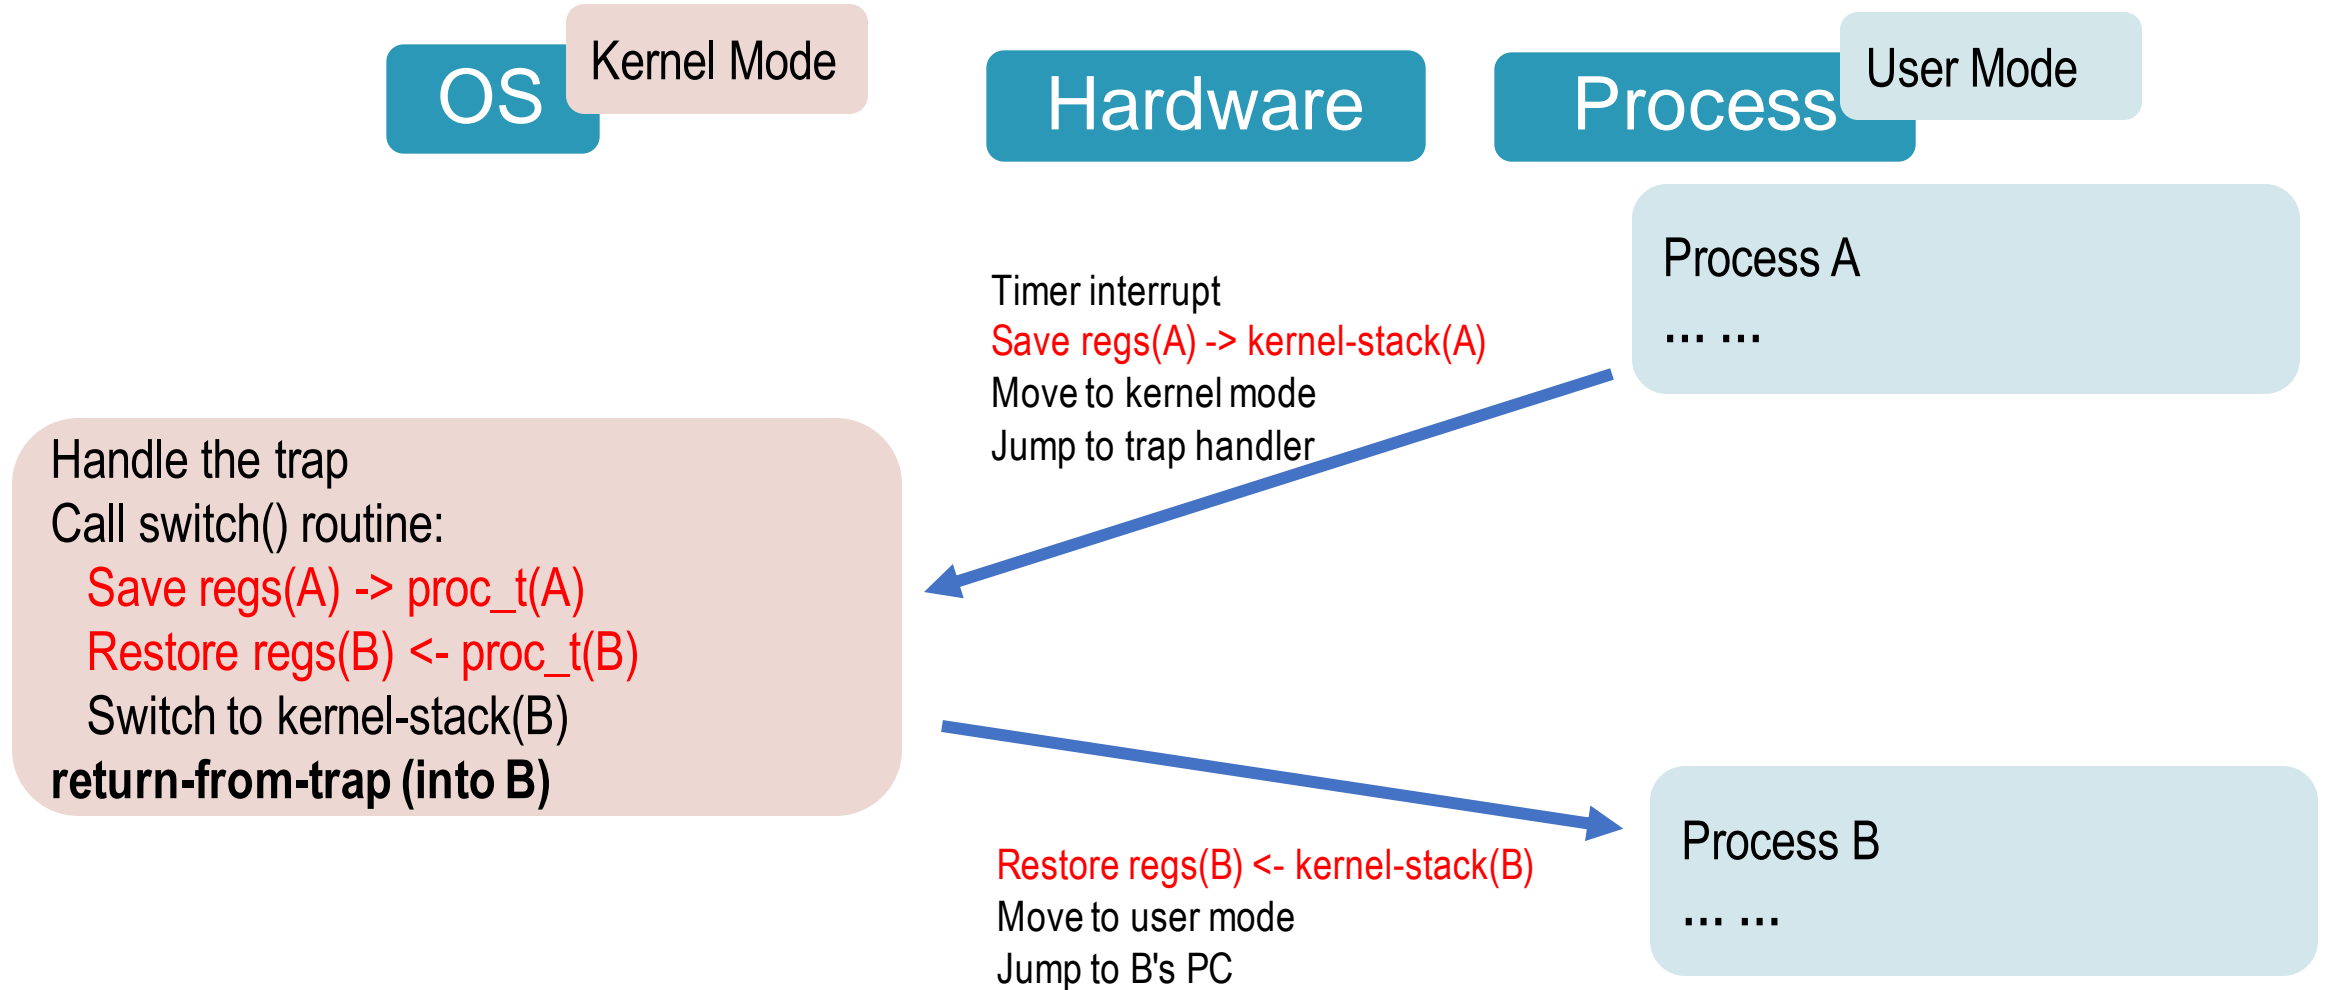
Remember address of syscall handler

Start interrupt timer

Start timer  
Interrupt CPU in X ms



# OS Regain Control via Timer Interrupt



# Process Control Block (struct proc) in xv6

```
proc.h

// Saved registers for kernel context switches.
struct context {
    uint64 ra;
    uint64 sp;

    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;
};

enum procstate { UNUSED, USED, SLEEPING,
RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum procstate state;           // Process state
    void *chan;                     // If non-zero, sleeping on channel
    int killed;                     // If non-zero, have been killed
    int xstate;                     // Exit status to be returned to parent's wait
    int pid;                        // Process ID

    // proc_tree_lock must be held when using this:
    struct proc *parent;            // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;                  // Virtual address of kernel stack
    uint64 sz;                      // Size of process memory (bytes)
    pagetable_t pagetable;          // User page table
    struct trapframe *trapframe;    // data page for trampoline.S
    struct context context;         // swch() here to run process
    struct file *ofile[NOFILE];    // Open files
    struct inode *cwd;              // Current directory
    char name[16];                  // Process name (debugging)
};
```

# switch.S

```

1  # Context switch
2  #
3  # void switch(struct context *old, struct context *new);
4  #
5  # Save current registers in old. Load from new.
6
7
8  .globl switch
9  switch:
10     sd ra, 0(a0)
11     sd sp, 8(a0)
12     sd s0, 16(a0)
13     sd s1, 24(a0)
14     sd s2, 32(a0)
15     sd s3, 40(a0)
16     sd s4, 48(a0)
17     sd s5, 56(a0)
18     sd s6, 64(a0)
19     sd s7, 72(a0)
20     sd s8, 80(a0)
21     sd s9, 88(a0)
22     sd s10, 96(a0)
23     sd s11, 104(a0)

```

store registers to old context  
sd = store doubleword command

```

24     ld ra, 0(a1)
25     ld sp, 8(a1)
26     ld s0, 16(a1)
27     ld s1, 24(a1)
28     ld s2, 32(a1)
29     ld s3, 40(a1)
30     ld s4, 48(a1)
31     ld s5, 56(a1)
32     ld s6, 64(a1)
33     ld s7, 72(a1)
34     ld s8, 80(a1)
35     ld s9, 88(a1)
36     ld s10, 96(a1)
37     ld s11, 104(a1)
38
39     ret

```

load new context to registers  
ld = load doubleword command

# Limited Direct Execution

*Time-sharing* is the idea that multiple processes can run together on a single machine as though they are in sole control of the machine

## How do processes share time?

*multiprogramming* – when a process waits for I/O, the OS can have another take over the CPU

*multitasking* – each process gets a time-slice, a time limit before the next process gets to execute on the CPU

## How does the OS keep control?

Hardware provides *interrupts*, *kernel mode* and *user mode*

System calls (cooperative) and interrupts (non-cooperative) enable OS to regain control after starting a process.