

# Recap

Direct Execution

Kernel/User modes => Restrict Direct Execution

- Context switch needed for mode switch

System call => Restricted User Process can attain its functional goals

Timer Interrupt => OS can regain control periodically and can switch processes

- Context switch needed for process switch

# Xv6 code flows

Kernel from beginning: entry.S

- > start() in start.c
- > main() in main.c
- > userinit() to start the first user process (user/init.c)
- > run scheduler() loop (in proc.c)

Handling trap:uservec in trampoline.S

- > usertrap() in trap.c
- > (for system call) syscall() in syscall.c -> ...->...
- > usertrapret() in trap.c
- > userret in trampoline.S

# L4: Scheduling

(Based on Chapter 7)

# Scheduling

The *workload* of a system consists of many processes with often different and conflicting requirements and user expectations

- Some processes just want a lot of CPU time (CPU bound)
- Others require a quick response to I/O (I/O bound)
- Users expect applications to be responsive (Interactive)

What **scheduling policies** will best meet user expectations?


# Jobs

Real processes alternate between needing CPU time and waiting for I/O

We call time when process needs CPU time a *job* (or *CPU burst*)

The time when process needs to wait is an *I/O burst*

Example: An OS that has only a single process looks like this



# Jobs With Time-Sharing

A scheduler may not let a process complete its job in one run on the CPU

On time-sharing systems, the scheduler may *preempt* a process, swapping it with a different process before it has completed its job

Example:

Process A has a CPU burst of 150ms

Process B has a CPU burst of 100ms




# Metrics

$T_{\text{arrival}}$  - time when job first enters ready state

$T_{\text{completion}}$  - time when job finishes

$T_{\text{firstrun}}$  - time when job starts its first run on the CPU

Example: Suppose jobs A and B both arrive at time 0 and execute as shown



A:  $T_{\text{firstrun}_A} = 0$ ,  $T_{\text{completion}_A} = 250$


B:  $T_{\text{firstrun}_B} = 50$ ,  $T_{\text{completion}_B} = 200$

# Metric: Turnaround Time

The time to complete a job

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

Example: Suppose jobs A and B both arrive at time 0 and execute as shown



$$A: T_{\text{turnaround}_A} = 250 - 0 = 250$$

$$B: T_{\text{turnaround}_B} = 200 - 0 = 200$$




# Metric: Response Time

The time to first execution on CPU

$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

Example: Suppose jobs A and B both arrive at time 0 and execute as shown



$$A: T_{\text{response}_A} = 0 - 0 = 0$$

$$B: T_{\text{response}_B} = 50 - 0 = 50$$

# Why Different Metrics?

Turnaround time tells time to complete jobs, good for *CPU bound* processes, where getting enough CPU runtime is the main concern

Response time tells how long to respond to I/O, good for *I/O bound (interactive)* processes, which have short CPU bursts and frequent I/O

There are other metrics, later we will look at *fairness*

# Policy: FIFO

Implementation: FIFO queue

Preemption: None

Advantage: Easy to implement

Example: Jobs arrive in the order A, B and C

	$T_{\text{arrival}}$	Runtime
A	0	50
B	0	50
C	0	200



$$T_{\text{average\_turnaround}} = (50 + 100 + 300) / 3 = 150$$

$$T_{\text{average\_response}} = (0 + 50 + 100) / 3 = 50$$

# Problem with FIFO

Example: Jobs arrive in the order A, B and C

	$T_{\text{arrival}}$	Runtime
A	0	200
B	0	50
C	0	50



$$T_{\text{average\_turnaround}} = (200 + 250 + 300) / 3 = 250$$

$$T_{\text{average\_response}} = (0 + 200 + 250) / 3 = 150$$

Performance depends on arrival order, large upfront CPU burst can hurt turnaround and response time

# Policy: Shortest Job First (SJF)


Implementation: Priority queue sorted by job length (shortest first)

Preemption: None

Advantage: **Optimal average turnaround time when all jobs arrive at same time**

Example: Jobs arrive in the order A, B and C

	$T_{\text{arrival}}$	Runtime
A	0	200
B	0	25
C	0	75



$$T_{\text{average\_turnaround}} = (25 + 100 + 300) / 3 = 142$$

$$T_{\text{average\_response}} = (0 + 25 + 100) / 3 = 42$$


# Problem with SJF

What happens if short jobs arrive after starting a long job?

Back to high turnaround and response times

Example:

	$T_{\text{arrival}}$	Runtime
A	0	200
B	5	25
C	10	75



$$T_{\text{average\_turnaround}} = ((200 - 0) + (225 - 5) + (300 - 10)) / 3 = 210$$

$$T_{\text{average\_response}} = ((0 - 0) + (200 - 5) + (225 - 10)) / 3 = 136.7$$

# Policy: Shortest Time-to-Completion First (STCF)


Implementation: Priority queue sorted by time to completion (shortest first)

Preemption: If new job arrives with a shorter time to completion, it preempts

Advantage: Short jobs don't need to wait for a long job to complete

Example: Jobs arrive in the order A, B and C

	$T_{\text{arrival}}$	Runtime
A	0	200
B	5	25
C	10	75



$$T_{\text{average\_turnaround}} = ((300-0) + (30-5) + (105-10)) / 3 = 140$$


$$T_{\text{average\_response}} = (0 + 0 + (30-10)) / 3 = 6.67$$

# Problem with STCF

Preemption gives STCF better response time in some cases, but there are still cases where response time can be poor

Example: Jobs arrive in the order A, B and C

	$T_{\text{arrival}}$	Runtime
A	0	200
B	0	200
C	0	200



$$T_{\text{average\_turnaround}} = (200 + 400 + 600) / 3 = 400$$

$$T_{\text{average\_response}} = (0 + 200 + 400) / 3 = 200$$



## Policy: Round Robin (RR)


## Implementation: FIFO queue

Preemption: Job on CPU gets *time-slice*, preempt when time expired

## Advantage: Low response time

Example: Jobs arrive in the order A, B and C; time-slice is 25

	$T_{\text{arrival}}$	Runtime
A	0	200
B	0	200
C	0	200



$$T_{\text{average\_turnaround}} = (550 + 575 + 600) / 3 = 575$$

$$T_{\text{average\_response}} = (0 + 25 + 50) / 3 = 25$$

# Pro and Con of RR


Pro: RR has great response time with guaranteed upper bound

$$\text{worst\_case\_response\_time} = \text{time\_slice} \times (\text{num\_jobs} - 1)$$

Con: Bad average turnaround time, frequent context switches reduce CPU efficiency

Example: Jobs arrive in the order A, B and C; time-slice is 25

	$T_{\text{arrival}}$	Runtime
A	0	200
B	0	200
C	0	200



$$T_{\text{average\_turnaround}} = (550 + 575 + 600) / 3 = 575$$

$$T_{\text{worst\_case\_response}} = 25 \times (3 - 1) = 50$$

# Another Problem with SJF and STCF: Oracle

Scheduler doesn't know how long it will take for a job to complete

SJF and STCF require an *oracle*, the ability to see into the future!

How can we get the benefits of STCF in a real system?