

Recap

Replacement Policies (for Physical Memory or TLB Cache)

- Optimal
- FIFO
- Random
- LRU

Clock Algorithm

LRU performs the closest in optimal for typical workloads, but it is really costly on every memory access

Clock Algorithm - is a way to approximate LRU cheaply

Add extra **use bit** to page table entry, on every memory access MMU set use bit of page to 1

When looking for page to evict

- visit pages in round-robin order
- if page use bit is 0, choose that page to evict
- else set use bit to 0 and continue search

Clock Algorithm: Example

Frames after accessing: 0 1 2 0 1

Frame #	0	1	2
Page	0	1	2
Use bit	1	1	1

No eviction; all pages are marked as used recently

To access page: 3

Frame #	0	1	2
Page	0 -> 3	1	2
Use bit	1 -> 0 -> 1	1 -> 0	1 -> 0

Page 0 (not exactly the LRU) is evicted; then page 3 (new) is marked as used recently

To access pages: 1 3

Frame #	0	1	2
Page	3	1	2
Use bit	1	0 -> 1	0

No eviction; pages 1 and 3 are marked as used recently

To access pages: 0

Frame #	0	1	2
Page	3	1	2 -> 0
Use bit	1	1 -> 0	0 -> 1

Page 2 (the LRU) is evicted; page 0 (new) is marked as used recently

Dirty Bit Optimization

If a page has only been read from (never written to) there is no reason to write it back to the swap space (if the swap page is large enough to contain a copy of it already) when it is evicted

Add a **dirty bit** to the page table, initialize to 0, on every write to the page table the MMU (hardware) sets the bit to 1


Can also make clock algorithm more efficient

- First try to find a page with dirty bit and use bit both 0, because it is lower cost to replace

What if All Accesses Were Random?

Without locality, none of the policies (except the cheating OPT) provide cost-effective benefit

If main memory is 50% of total pages, only 50% hit rate



80-20 Workload


More typical, the **80-20 workload** (80% of accesses are to 20% of pages)

A few pages get most accesses

Most pages get few accesses

Results from locality (either spatial or temporal)

LRU performs closest to optimal, FIFO and RAND perform the same

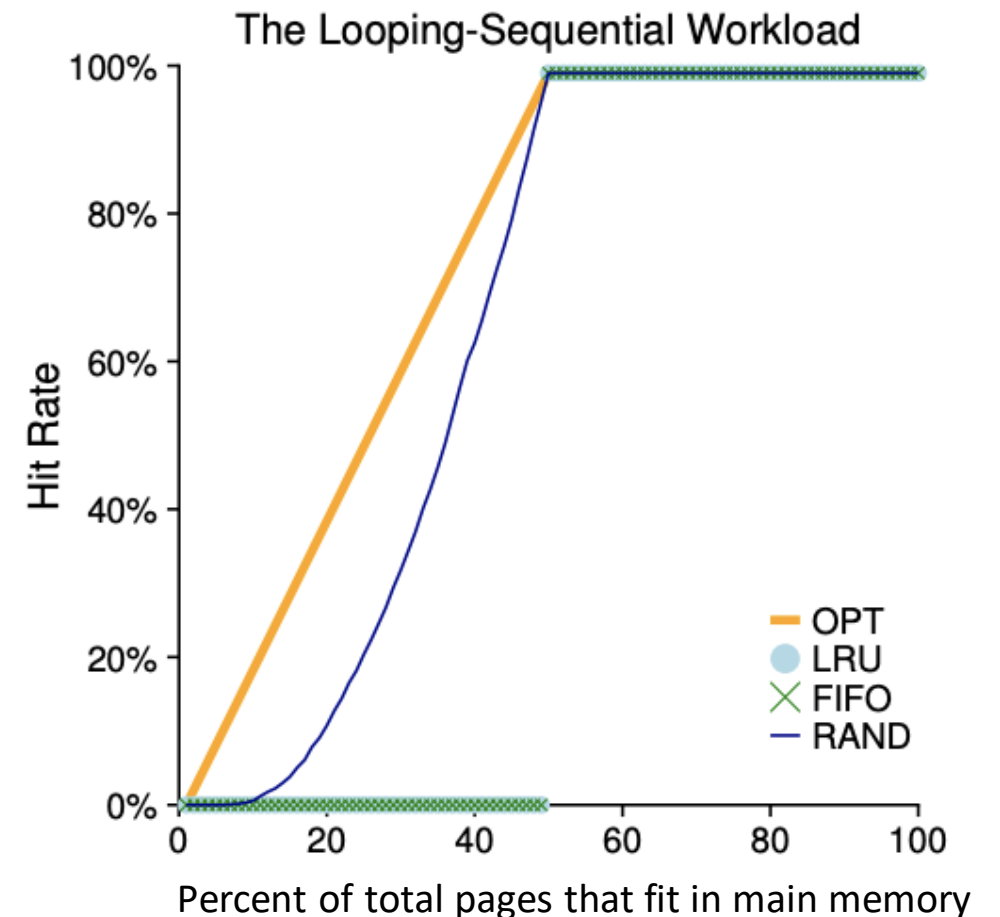


Looping-Sequential Workload

Assume loop repeatedly reads pages 0 to 49 in increasing order


When cache size is below 50, LRU and FIFO have same corner case that causes every access to be a miss

Random avoids corner cases



Clock Policy with 80-20 Workload

Clock is a very close approximation of LRU



Thrashing

Example: System has two processes that both sequentially read from N pages in a continuous loop. The System only has enough main memory to store N pages.

Thrashing is when process is spending more time on handling page-fault than executing

- Starts at one processes and snowballs into several processes thrashing
- Multiprogramming and multitasking make it worse, not better
- Sudden and extreme drop in system performance

Need to reduce the amount of multiprogramming and multitasking to avoid thrashing

L13: Concurrency

(based on Ch. 26)

Concurrency

We have explored how OSes use concurrency

- multiprogramming – to utilize the CPU efficiently when programs block for I/O
- multitasking – to give all programs a fair slice of the CPU and make progress

What about providing concurrency within a process?

- Programs have tasks that need to block for I/O
- Programs want to be responsive to the user while performing tasks in the background
- Programs want to distribute their computations across multiple CPUs to complete faster

Processes can have multiple **threads** - concurrent points of execution

How does the OS provide programs concurrent threads of execution?

What are Threads?

Multiple
points of
execution
of a
program
within a
single
process

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  static volatile int counter = 0;
7
8  // mythread()
9  //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }
```

Thread2 →

Thread1 →

Thread0 →

What are Threads?

Multiple points of execution means each thread must have its own program counter (PC - next instruction to execute) and call stack

Thread2
PC
Stack2

Thread1
PC
Stack1

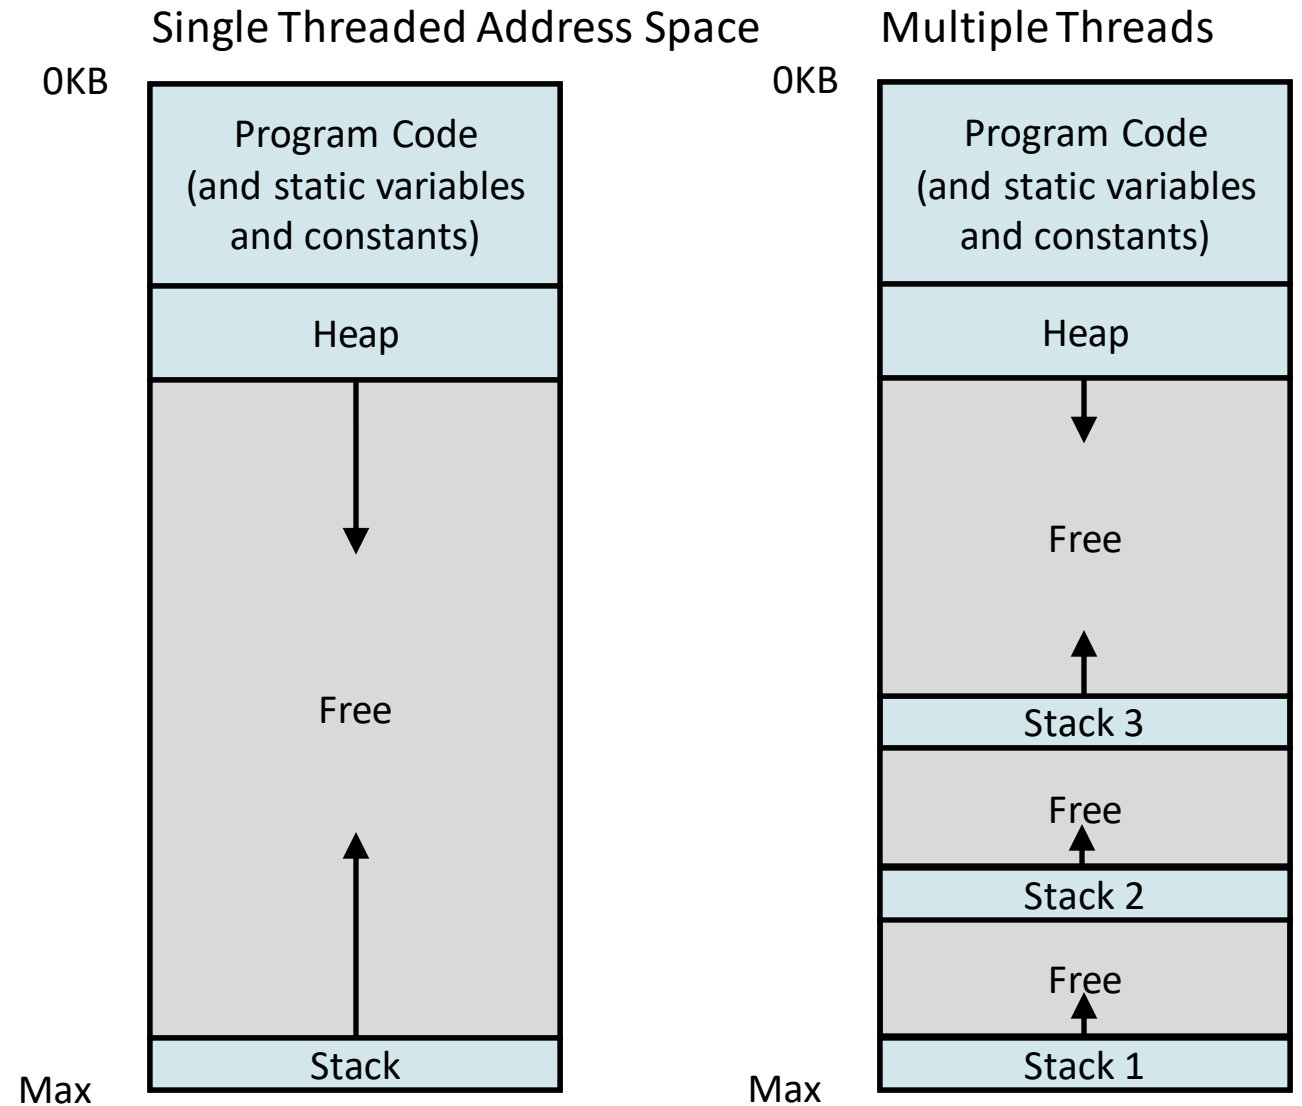
Thread0
PC
Stack0

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  static volatile int counter = 0;
7
8  // mythread()
9  //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
23
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39           counter);
40     return 0;
41 }
```

Multi-Threaded Address Space

Each thread has its own stack segment

Program, data and heap are shared with the process



Thread Control Block (TCB)

Recall the Process Control Block (PCB)

- Keeps track of information for each process
- Stores execution context to enable context switching out of and back into the process

Thread Control Block (TCB) is similar, but

- threads share process code, data and heap segments

Scheduler uses both PCBs and TCBs to decide which thread to run next

Process Control Block

Process ID (pid)
State (e.g., running, runnable, blocked)
Program Counter
CPU Register values (context)
Stack pointer

Pointers to code, data and heap segments
Pointer to PCB of the parent process
Open file descriptors

Thread Control Block

Thread ID (tid)
State (e.g., running, runnable, blocked)
Program Counter
CPU Register values (context)
Stack pointer

Pointer to the PCB of the process that thread belongs to

Concurrency vs Parallelism

Concurrent means multiple threads making progress in time but may be implemented by time-sharing, can be on one or more CPU cores

Parallel means multiple threads instructions executing independently on multiple CPU cores



Concurrency of threads enables

I/O overlap - can overlap blocking I/O with other program tasks (same concept as multiprogramming)

Responsiveness - user can continue to interact with system even when program is performing heavy processing in the background

Parallelism of threads enables

Performance – finish more tasks in less time by distributing load to multiple CPU cores



Are Threads Needed?

What about multiple processes?

- Xv6 does not have user threads

- Can use `fork()`, `pipe()` and `wait()` to manage concurrent processes

Threads provide more convenience and better performance

- Simple memory sharing (all threads share same data and heap)

- Lower cost of thread creation (don't need to allocate new address space, just stack)

- Lower cost of context switch (only stack and registers change)

Example

```
#include <stdio.h>
#include <pthread.h>

static volatile int counter = 0;

void *mythread(void *arg) {
    printf("thread %s: begin\n", (char *) arg);
    for (int i = 0; i < 1e7; i++) {
        counter++;
    }
    printf("thread %s: end\n", (char *) arg);
    return NULL;
}

int main() {
    pthread_t p1, p2;
    printf("main: begin\n");

    pthread_create(&p1, NULL, mythread, "A");
    pthread_create(&p2, NULL, mythread, "B");

    pthread_join(p1, NULL);
    pthread_join(p2, NULL);

    printf("main: done with both (counter = %d)\n", counter);
    return 0;
}
```

When can Context Switch Occur?

Scheduler can decide to context switch to another thread at any instruction

High level code

```
counter = counter + 1;
```

Assembly instructions

```
mov 0x8049a1c, %eax ← Can switch here  
add $0x1, %eax ← Or here  
mov %eax, 0x8049a1c ← Or here
```

The Problem (Race Condition)

Concurrent update of shared memory can result in **race condition** bug

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	<i>before critical section</i>		100	0	50
	mov 8049a1c,%eax		105	50	50
	add \$0x1,%eax		108	51	50
interrupt					
<i>save T1</i>					
<i>restore T2</i>			100	0	50
		mov 8049a1c,%eax	105	50	50
		add \$0x1,%eax	108	51	50
		mov %eax,8049a1c	113	51	51
interrupt					
<i>save T2</i>					
<i>restore T1</i>			108	51	51
	mov %eax,8049a1c		113	51	51