

Recap

File System Interface

- File & Directory
- POSIX API: open/close, read/write, random access, hard/soft link

File System Implementation

Based on Ch. 40

File System Implementation

Many possible approaches to making a file system

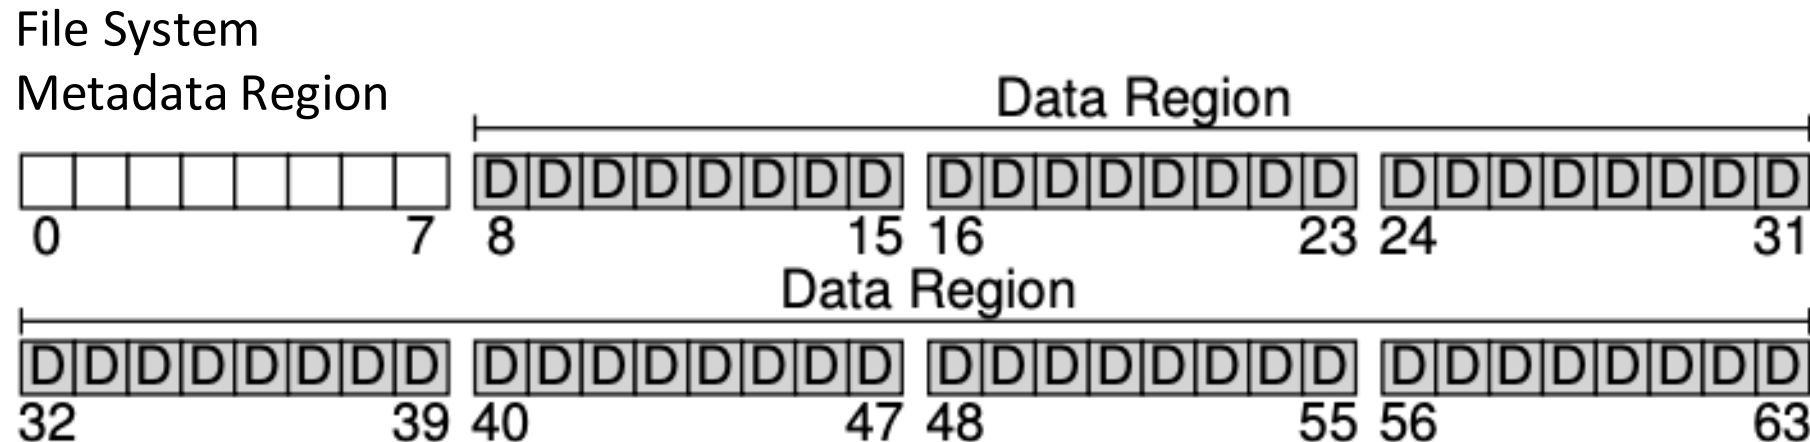
Need to understand the **data structures** for files, directories and free space and how **access** (e.g., read and write) is implemented

How to implement a simple file system?

A Very Simple File System (VSFS)

Disk is divided into fixed sized blocks (e.g., block size = 4KB)

A few blocks are reserved for file system metadata (information about the files), the rest stores the data of the files



Contiguous vs Block Allocation

With memory we explored base and bounds (contiguous allocation) or paging


File system have a similar choice (blocks are like pages)

Contiguous allocation has problem of **external fragmentation**

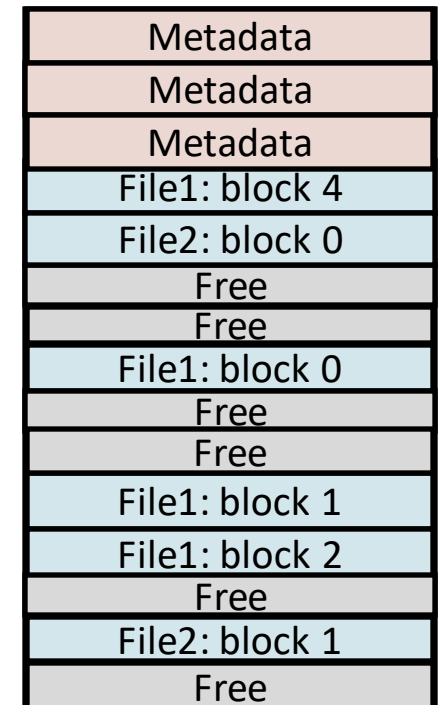
Block allocation has problem of **internal fragmentation**

Block allocation usually wins because internal fragmentation is less of a problem and block allocation is much more efficient in allocating and resizing files

Contiguous



Blocks



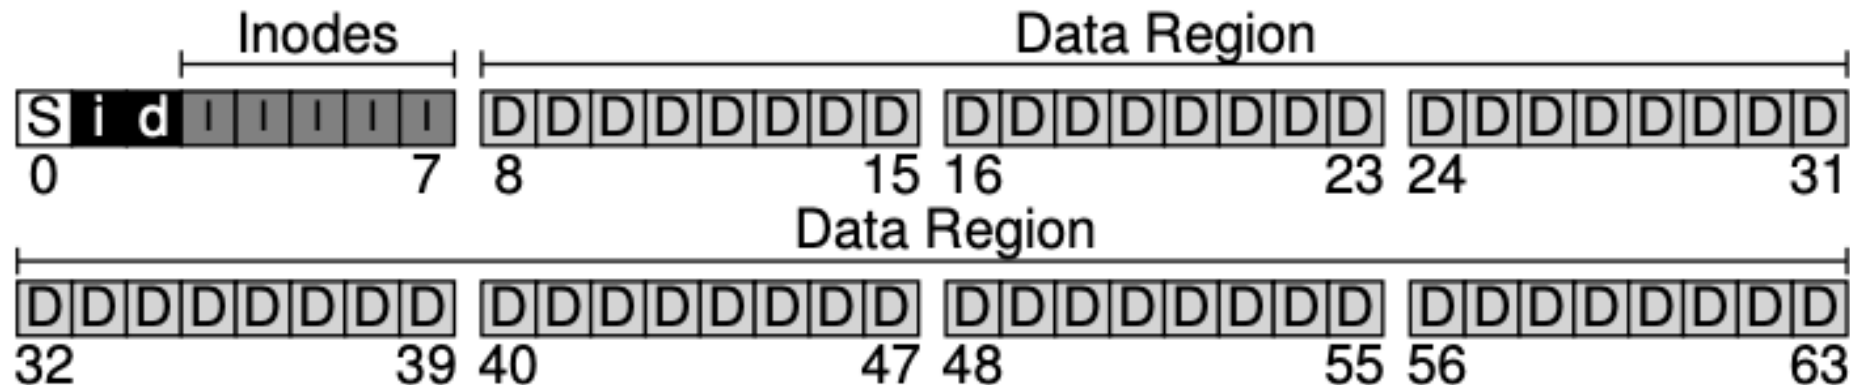
File System Metadata

inodes (I) contains information about a particular file

data bitmap (d) stores which blocks are free/used in the data region

inode bitmap (i) stores which blocks are free/used in the metadata region

Super block (S) contains information about the file system



inode

inode (index node) contains metadata (information) about a file

Example: Ext2 file system inode

Size	Name	What is this inode field for?
2	mode	can this file be read/written/executed?
2	uid	who owns this file?
4	size	how many bytes are in this file?
4	time	what time was this file last accessed?
4	ctime	what time was this file created?
4	mtime	what time was this file last modified?
4	dtime	what time was this inode deleted?
2	gid	which group does this file belong to?
2	links_count	how many hard links are there to this file?
4	blocks	how many blocks have been allocated to this file?
4	flags	how should ext2 use this inode?
4	osd1	an OS-dependent field
60	block	a set of disk pointers (15 total)
4	generation	file version (used by NFS)
4	file_acl	a new permissions model beyond mode bits
4	dir_acl	called access control lists

← points to other inode or data blocks

Direct Indexing

In **direct indexing** the inode for a file has pointers to the data blocks of the file

Suppose an inode has 15 block pointers and blocks are 4KB, what is the largest possible file size?

$$15 * 4KB = 60KB \text{ maximum file size}$$

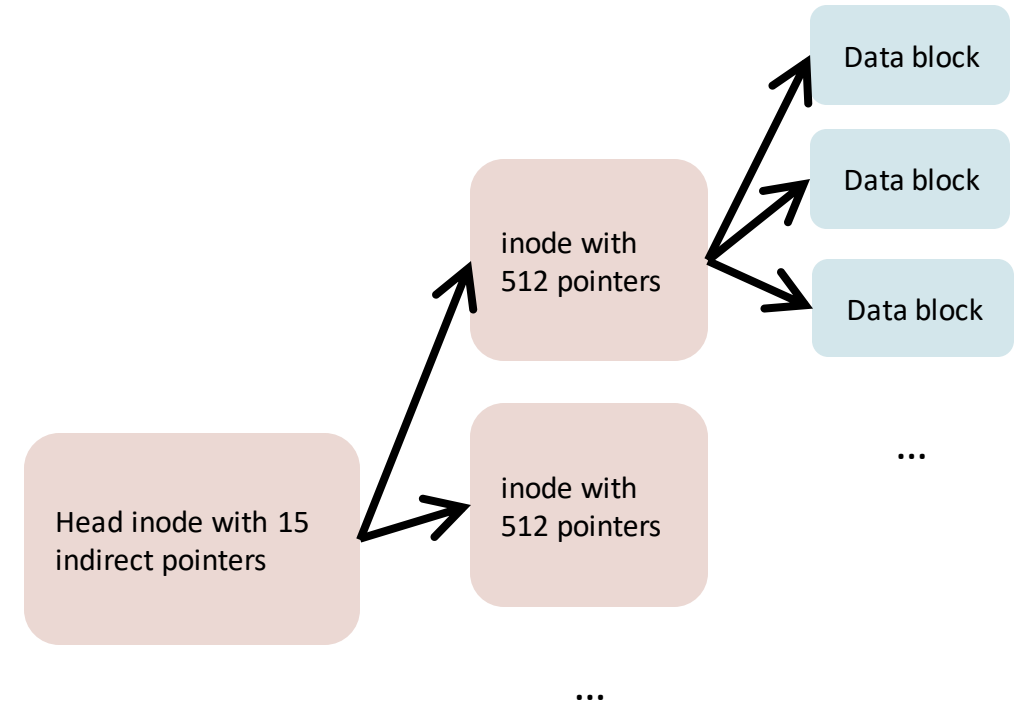
Indirect indexing

An **indirect pointer** points to an inode with more pointers

Suppose pointers are 8 bytes and block size is 4KB, then a block can hold 512 pointers


With one level of indirection the previous example has:

$15 * 512 * 4\text{KB} = \text{about } 30\text{MB}$ maximum file size



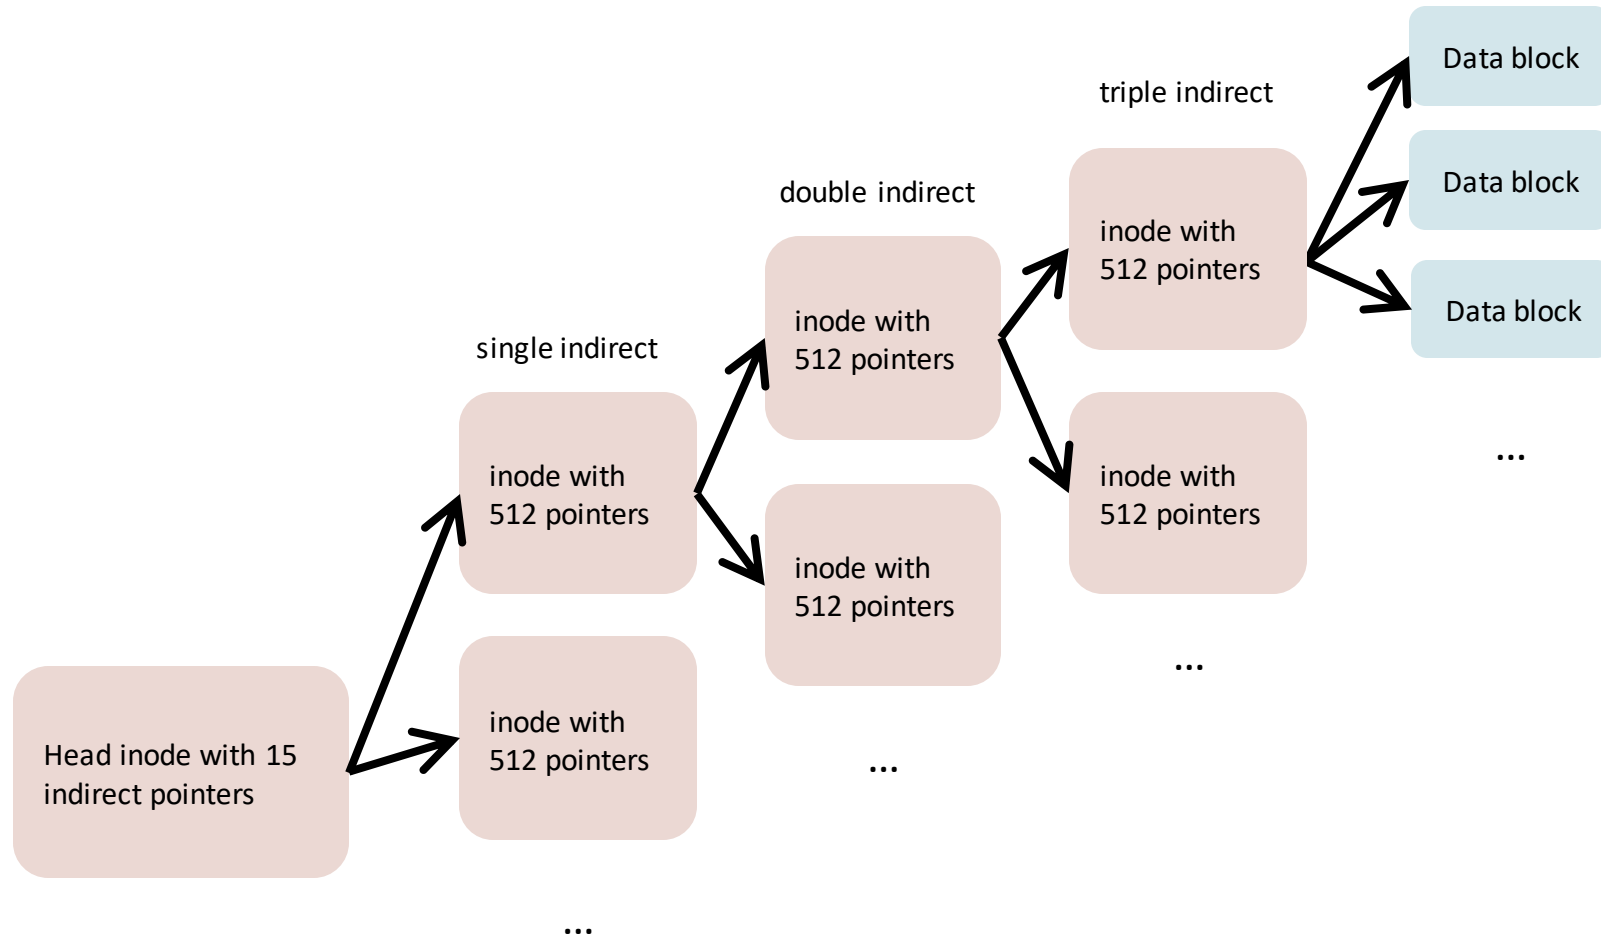
Multiple Levels of Indirection

Double indirect means two levels of indirection



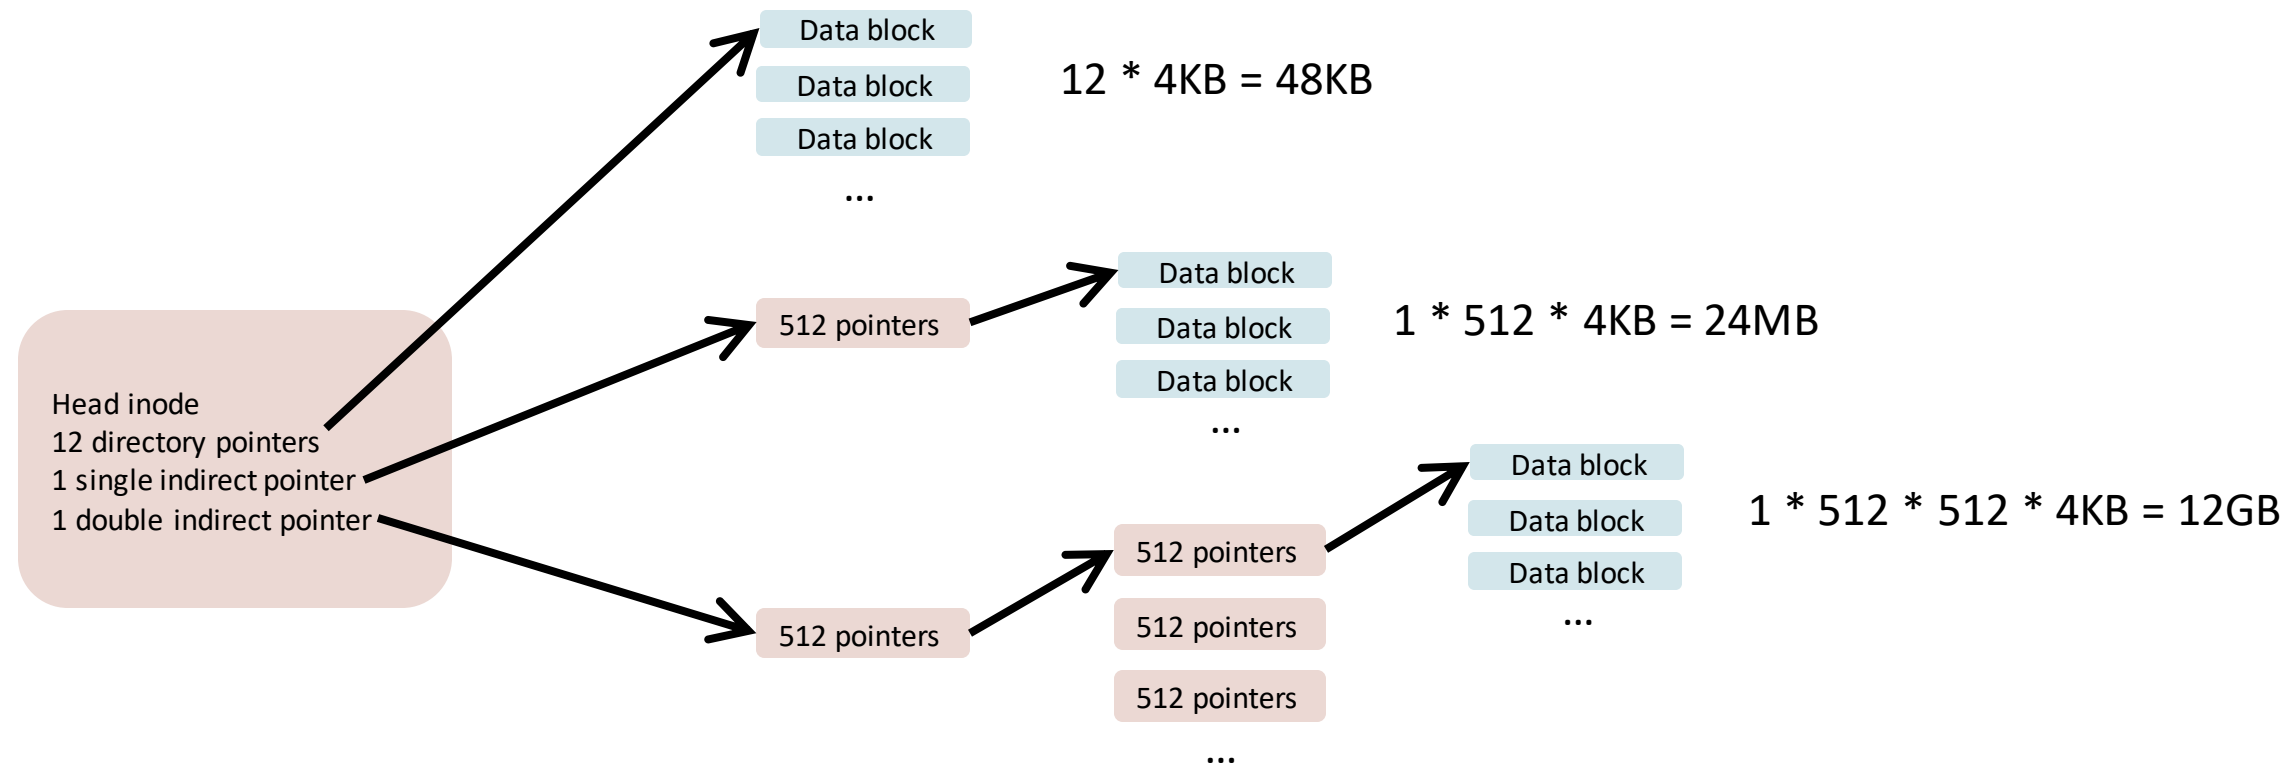
Multiple Levels of Indirection

Triple indirect means three levels of indirection



Multi-Level Indexing

Head inode may combine multiple levels of indirection



Common Observations of File Systems

Most files are small	~2K is the most common size
Average file size is growing	Almost 200K is the average
Most bytes are stored in large files	A few big files use most of space
File systems contains lots of files	Almost 100K on average
File systems are roughly half full	Even as disks grow, file systems remain ~50% full
Directories are typically small	Many have few entries; most have 20 or fewer

Example Reading a File

How to read from the file /foo/bar? **First,**
fd=open("/foo/bar").

1. inode number for root directory is well known; read "root inode" and get the location for "root data".
2. Read "root data" to get inode number for "/foo".
3. Read "foo inode" to get location of "foo data".
4. Read "foo data" to get inode number for "/foo/bar".
5. Read "bar inode" and bring the metadata of /foo/bar to main memory.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
open(bar)			read	read		read	read			
read()					read			read		
read()					write read				read	
read()					write read					read
					write					

Example Reading a File

How to read from the file /foo/bar?

Second, read(fd,...).

1. Read "bar inode" to get the location of the first data block of /foo/bar, i.e., "bar data [0]".
2. Read "bar data [0]".
3. Update last access time at "bar inode".
4. Read "bar inode" to get the location of the second block of /foo/bar, i.e., "bar data[1]".
5. Read "bar data [1]".
6. Update last access time at "bar inode".
7.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
open(bar)			read	read		read	read			
read()					read			read		
read()					write					
read()					read				read	
read()					write					
					read					read
					write					

Example: Creating and Writing File

How to create file /foo/bar?

1. Read "root inode" (i.e., inode for "/") to get location of "/".
2. Read data block of "/" (i.e., "root data") to get inode number of "/foo".
3. Read "foo inode" to get location of "/foo".
4. Read "/foo" (i.e., "foo data"); read inode bitmap to find an empty inode, denoted as x, and update the bitmap; add a pair ("bar", x) to "foo data".
5. Read "bar inode" (i.e., inode with number x) and initialize it.
6. Update last access time at "foo inode".

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
create (/foo/bar)		read write	read	read		read		read		
write()	read write			write	read write			write		
write()	read write				read					
write()	read write				write			write		
write()	read write				read					write

Example: Creating and Writing File

How to **write** to new file /foo/bar?

1. Read "bar inode".
2. Read data bitmap to find an empty data block x and update the bitmap.
3. Write to data block x (i.e., "bar data [0]").
4. Add block x to "bar inode".
5. Update the last access time at "bar inode"
6. Read "bar inode".
7. Read data bitmap to find an empty data block y and update the bitmap.
8. Write to data block y (i.e., "bar data [1]").
9. Add block y to "bar inode".
10.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
create (/foo/bar)		read write	read	read		read	read	write		
write()	read write				read write				write	
write()	read write				write read				write	
write()	read write				write read					write

Basic Performance Improvements

Caching – holds popular blocks to decrease number of times blocks are read from disk

Write buffering - batch multiple updates into a smaller set of I/O operations