

Recap

File System Implementation – Data Structures

- Data blocks (fixed size)
- Metadata: Superblock, inode bitmap, data bitmap, inodes
 - Data block Indexes in inodes: direct pointers; single/double/triple indirect pointers

Example Reading a File

How to read from the file /foo/bar? **First,**
fd=open("/foo/bar").

1. inode number for root directory is well known; read "root inode" and get the location for "root data".
2. Read "root data" to get inode number for "/foo".
3. Read "foo inode" to get location of "foo data".
4. Read "foo data" to get inode number for "/foo/bar".
5. Read "bar inode" and bring the metadata of /foo/bar to main memory.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
open(bar)			read	read		read	read			
read()					read			read		
read()					write read				read	
read()					write read					read
					write					

Example Reading a File

How to read from the file /foo/bar?

Second, `read(fd,...)`.

1. Read "bar inode" to get the location of the first data block of /foo/bar, i.e., "bar data [0]".
2. Read "bar data [0]".
3. Update last access time at "bar inode".
4. Read "bar inode" to get the location of the second block of /foo/bar, i.e., "bar data[1]".
5. Read "bar data [1]".
6. Update last access time at "bar inode".
7.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
open(bar)			read	read		read	read			
read()					read			read		
read()					write					
read()					read				read	
read()					write					
					read					read
					write					

Example: Creating and Writing File

How to **create** file /foo/bar?

1. Read "root inode" (i.e., inode for "/") to get location of "/".
2. Read data block of "/" (i.e., "root data") to get inode number of "/foo".
3. Read "foo inode" to get location of "/foo".
4. Read "/foo" (i.e., "foo data"); read inode bitmap to find an empty inode, denoted as x, and update the bitmap; add a pair ("bar", x) to "foo data".
5. Read "bar inode" (i.e., inode with number x) and initialize it.
6. Update last access time at "foo inode".

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
create (/foo/bar)		read write	read	read		read	read write			
write()	read write			write	read write					
write()	read write				read		write			
write()	read write				write read			write		
write()	read write				write read					write

Example: Creating and Writing File

How to **write** to new file /foo/bar?

1. Read "bar inode".
2. Read data bitmap to find an empty data block x and update the bitmap.
3. Write to data block x (i.e., "bar data [0]").
4. Add block x to "bar inode".
5. Update the last access time at "bar inode"
6. Read "bar inode".
7. Read data bitmap to find an empty data block y and update the bitmap.
8. Write to data block y (i.e., "bar data [1]").
9. Add block y to "bar inode".
10.

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]	bar data [2]
create (/foo/bar)		read write	read	read		read	read	write		
write()	read write				read write				write	
write()	read write				write read				write	
write()	read write				write read					write

Basic Performance Improvements

Caching – holds popular blocks to decrease number of times blocks are read from disk

Write buffering - batch multiple updates into a smaller set of I/O operations

Fast File System

Based on Ch. 41

Faster File System

Original UNIX file system had very poor performance (used only 2% of disk bandwidth)

The most significant physical limitation is the difference between random and sequential latencies


Optimizations need to take the disk into account

How to improve file system performance?

Consequence of Fragmentation

Block based allocation means files can become spread out over the disk
Best performance is when files are written to contiguous memory


1. Assume files A, B, C and D are stored on disk



2. B and D are deleted leaving two gaps




3. Blocks of E are spread out over the disk



Cylinder Groups


A **cylinder** is a set of tracks near to each other on the drive



Block Groups

Because most hard drives don't provide enough information to choose cylinder groups, most file systems are organized by block groups


Block groups are consecutive portions of the disk's address space



The Berkeley Fast File System (FFS)

Principle: *keep related stuff together*

A single block group (file system has many)



Each block group has a copy of superblock (S), inode blocks and data blocks


Two heuristics to improve performance:

- Try to allocate data blocks for a file in the same block group as the file's inode
- Try to locate files that are together in a directory in the same block group

Path Locality

FFS heuristics are based on another form of locality

Path Locality – consecutive file accesses are likely to be to file paths that are near to each other



Large File Exception

Large file will completely fill block group, preventing files in same directory being in same group

Heuristic

After blocks are allocated into the first block group (e.g., the 12 direct pointers), FFS places the next “large chunk” (e.g., those pointed to by the first indirect block) in another block group

Sub-Block

Block Size is 4KB, but most of small files have size $< 2K$. Internal Fragmentation!

For small files, allocate smaller blocks (512KB) called sub-blocks.

Writing small blocks is inefficient (more frequent positioning); for better efficiency, buffer writing till large block(s) of data to write.